

An ensemble model for the machine reading comprehension dataset SQuAD

Summer report

Maria Fernanda Mora

September 12th, 2016



Carnegie Mellon University
Language Technologies Institute

IT&M

Overview

- Problem definition
- Exploratory analysis
- Pipeline description
- Sentence ranking
- Answer extraction
- Implementation

Problem definition

Problem definition

- Implement a system capable of performing reading comprehension over SQuAD's data set that outperforms the current state of the art.
- *SQuAD's challenge:*
 - No candidate answers provided
 - A correct answer to a question can be any sequence of tokens from the given text
 - Q&A in SQuAD were created by humans, hence more realistic

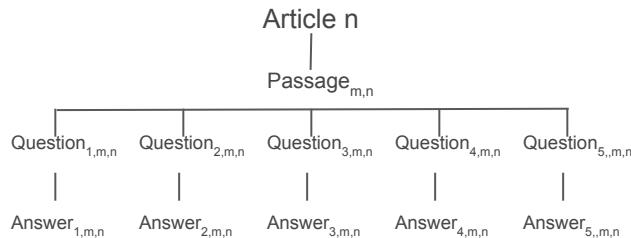
Exploratory analysis

Exploratory analysis

- General statistics
- Lexical analysis
- Syntactic analysis

Complete dataset

- 536 Wikipedia articles
- 108K QA pairs
- Training, dev and test
- Hierarchical view:



Model evaluation

- Output: sequence of tokens
- Measures: Exact match, F1

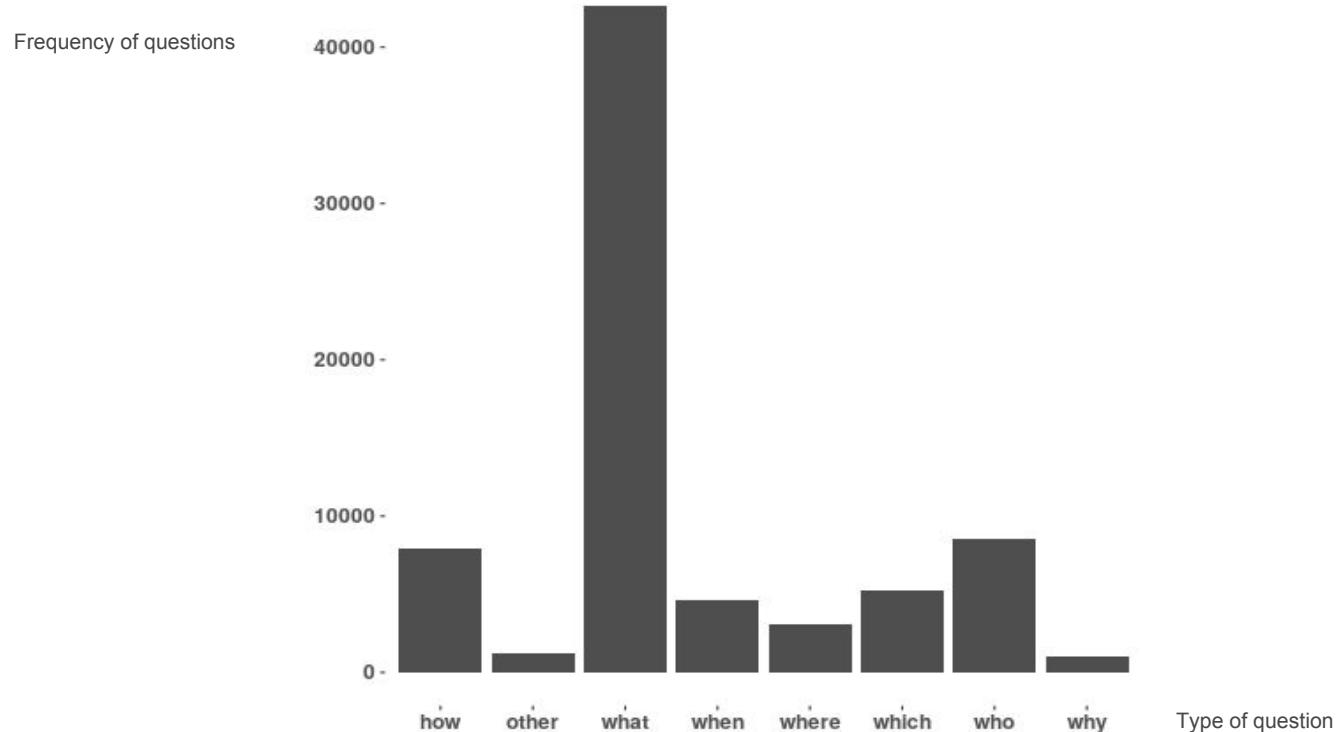
Training dataset

- 378 Wikipedia articles
- ~ 42 passages per article
- 5 questions per passage
- 1 answer per question
- ~ 80K QA pairs

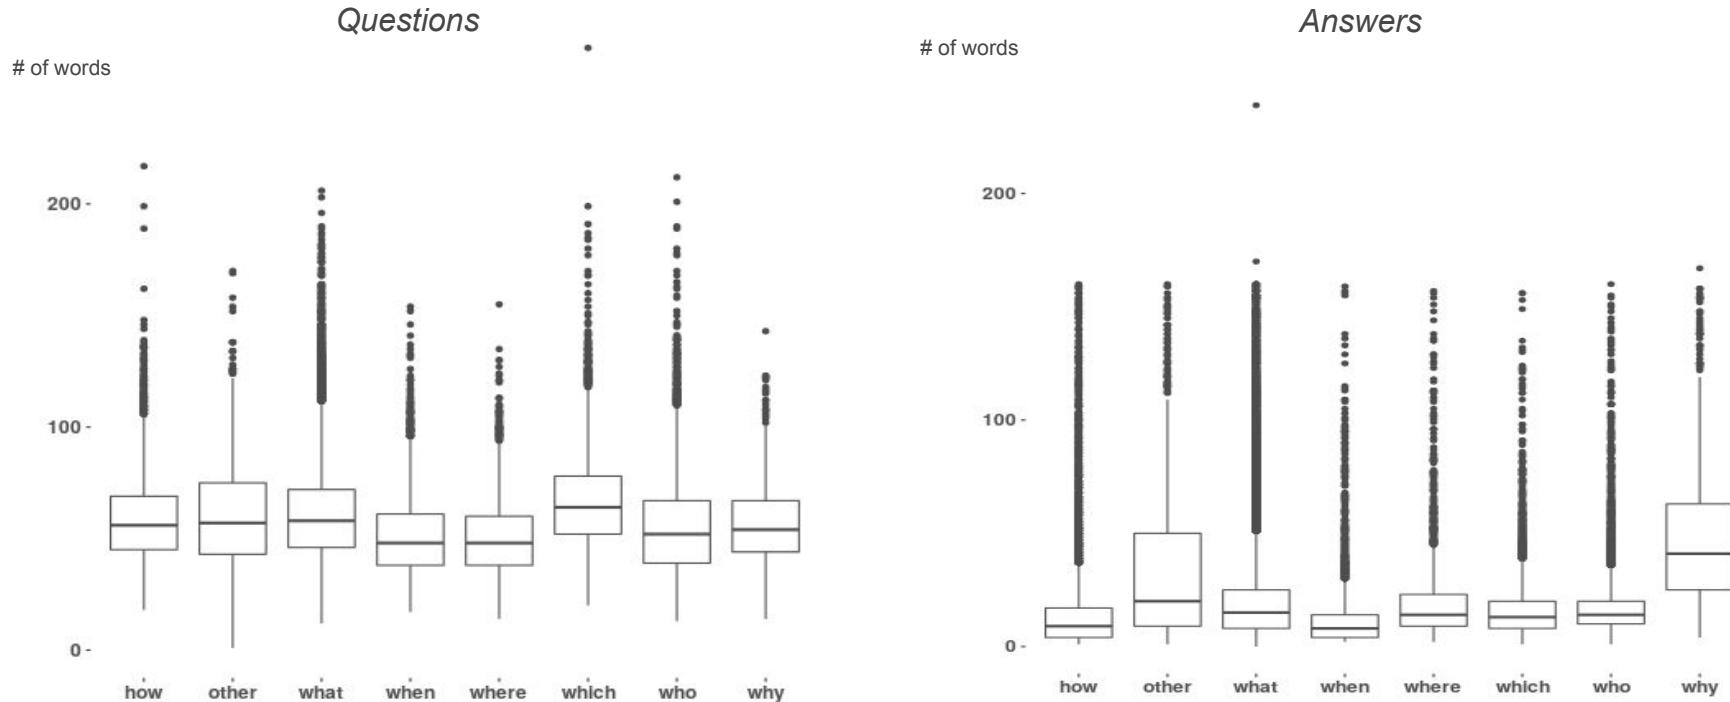
Vocabulary Size

	# words
Passages	~88K (98% without stop words)
Questions	~1K (93% w/o stop words)
Answers	~0.5K (93% w/o stop words)

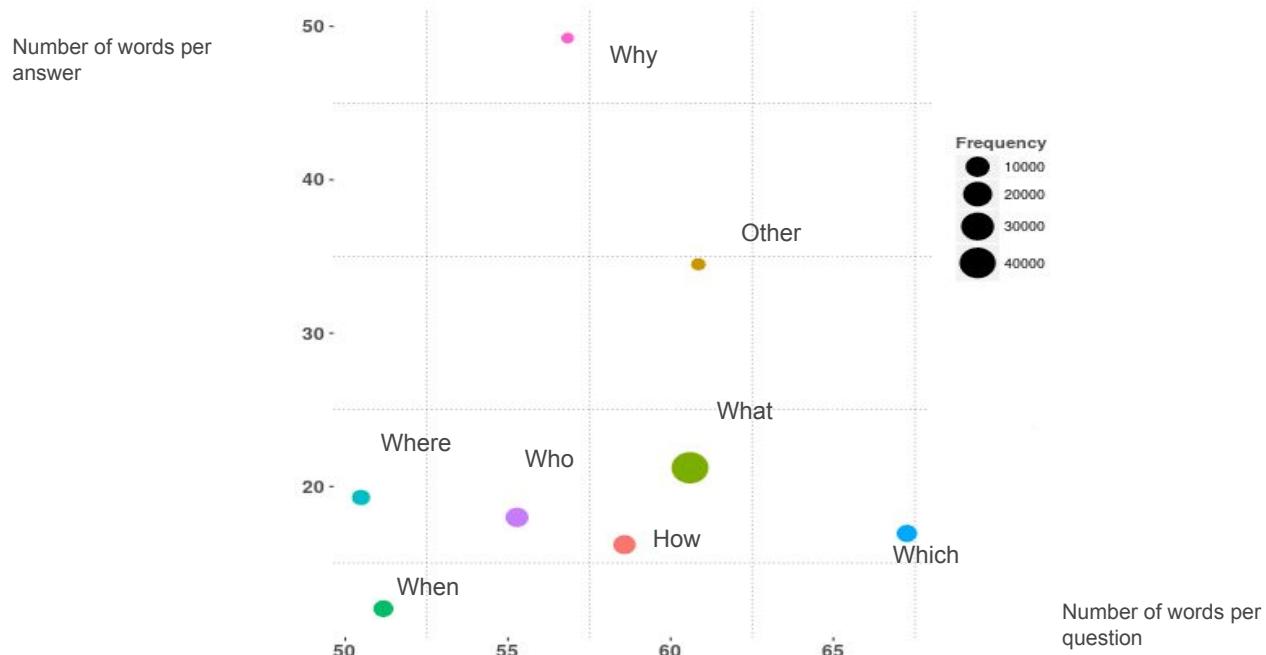
>99% of the questions are factoid; >50% are *what* questions



Questions length is similar; answers to *why* and *other* questions show length variation



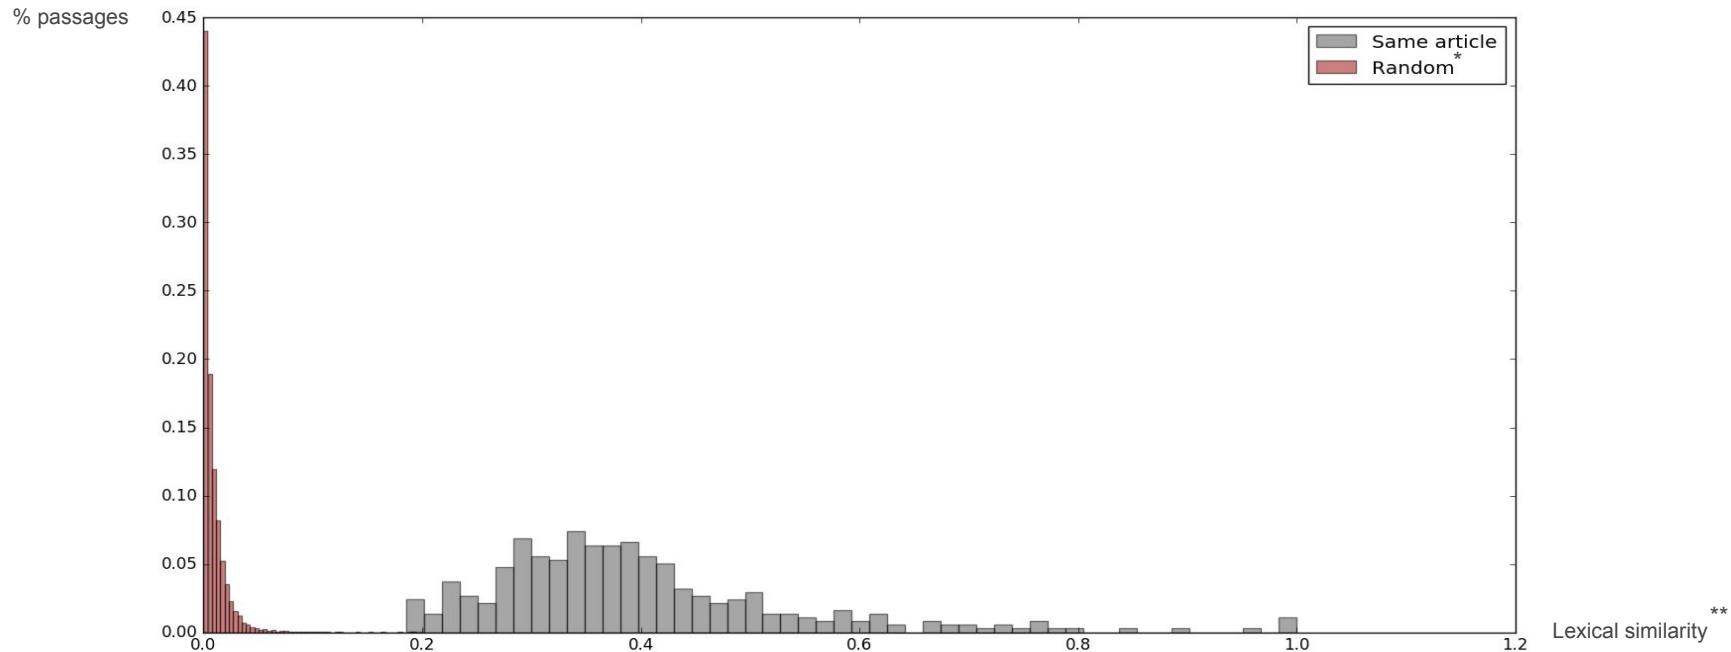
Questions are larger than answers; *why* questions have the largest answers but represent <5%



Exploratory analysis

- General statistics
- Lexical analysis
- Syntactic analysis

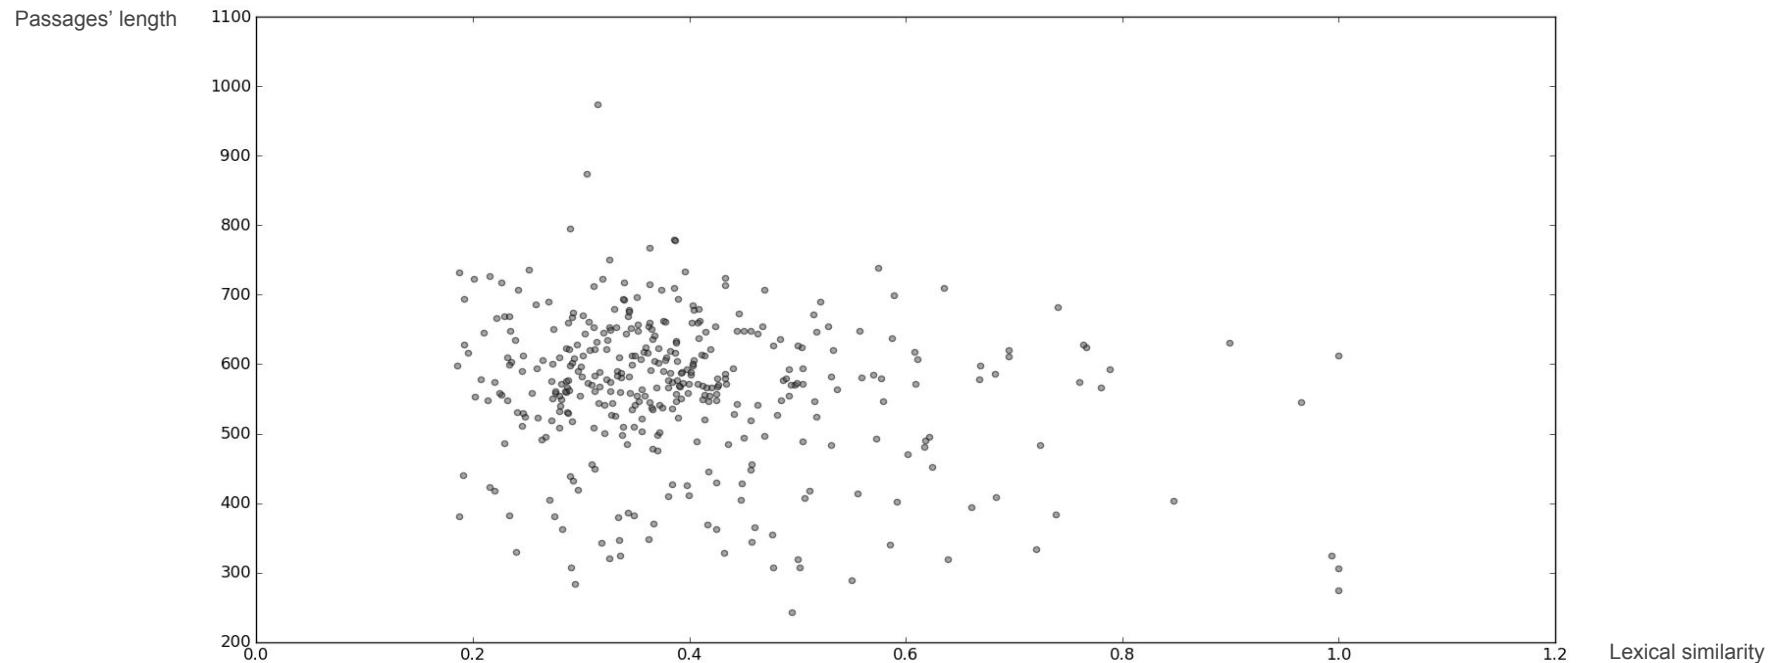
There exists a lexical similarity 0.3-0.4 between passages of the same article



* Random passages were extracted from all the articles

** Measured as cosine similarity

This similarity is independent of the length of the passage



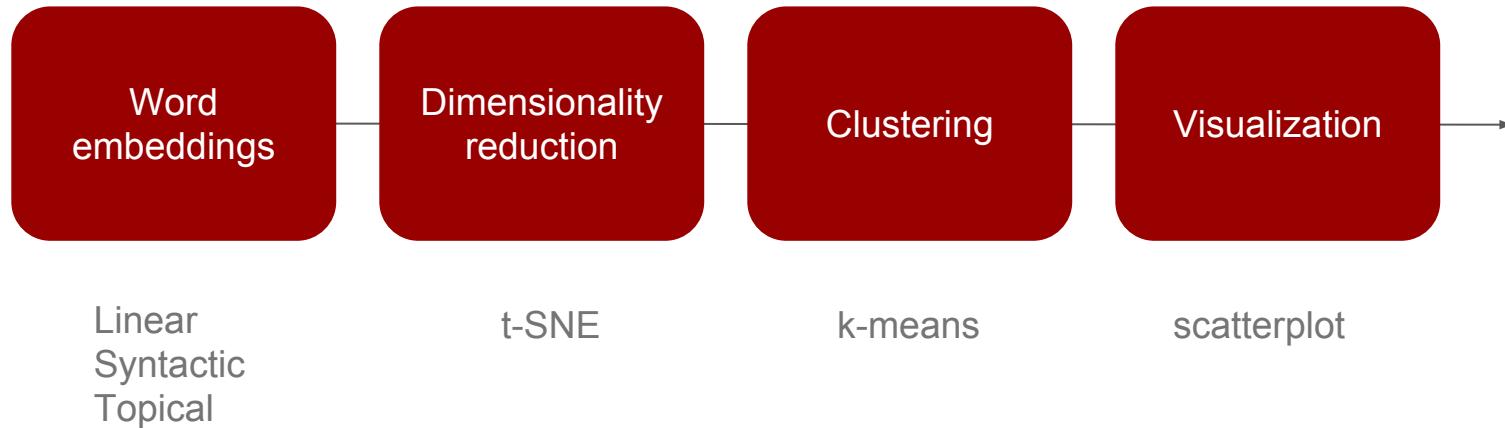
LDA analysis varying number of words and topics showed the following persistent topics

- history
- government
- nation-state
- sports
- art

Exploratory analysis

- General statistics
- Lexical analysis
- Syntactic analysis
 - Embeddings
 - Word
 - Sentence
 - Paragraph

Word embeddings pipeline



Models

- Glove
- Skip-gram

Parameters

- Window size
- Vector size

GLOVE

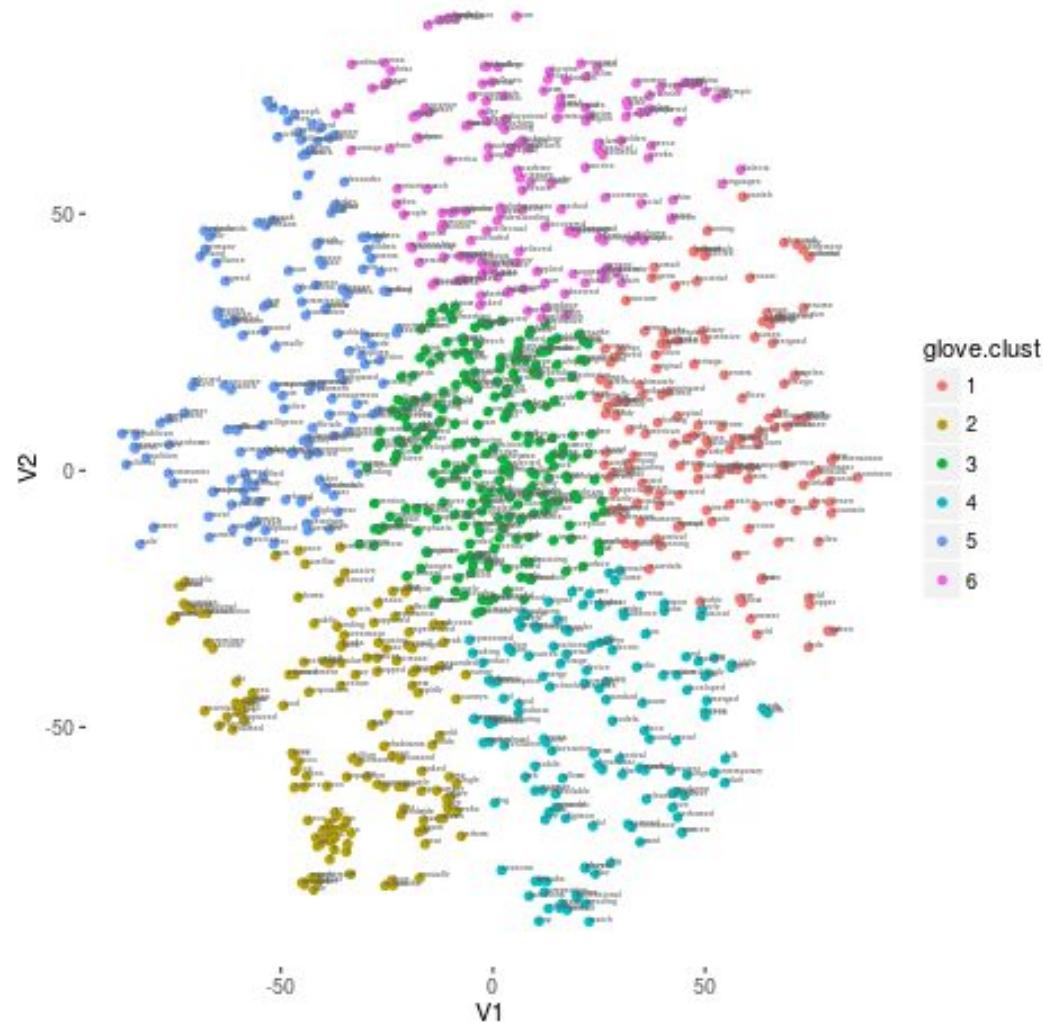
Linear Embedding

- Min words in voc = 100
- Size of vectors = 100, 300, 500
- Size of window = 5, 15, 20

GLOVE Linear Embedding

Window Size = 15

Vector Size = 100

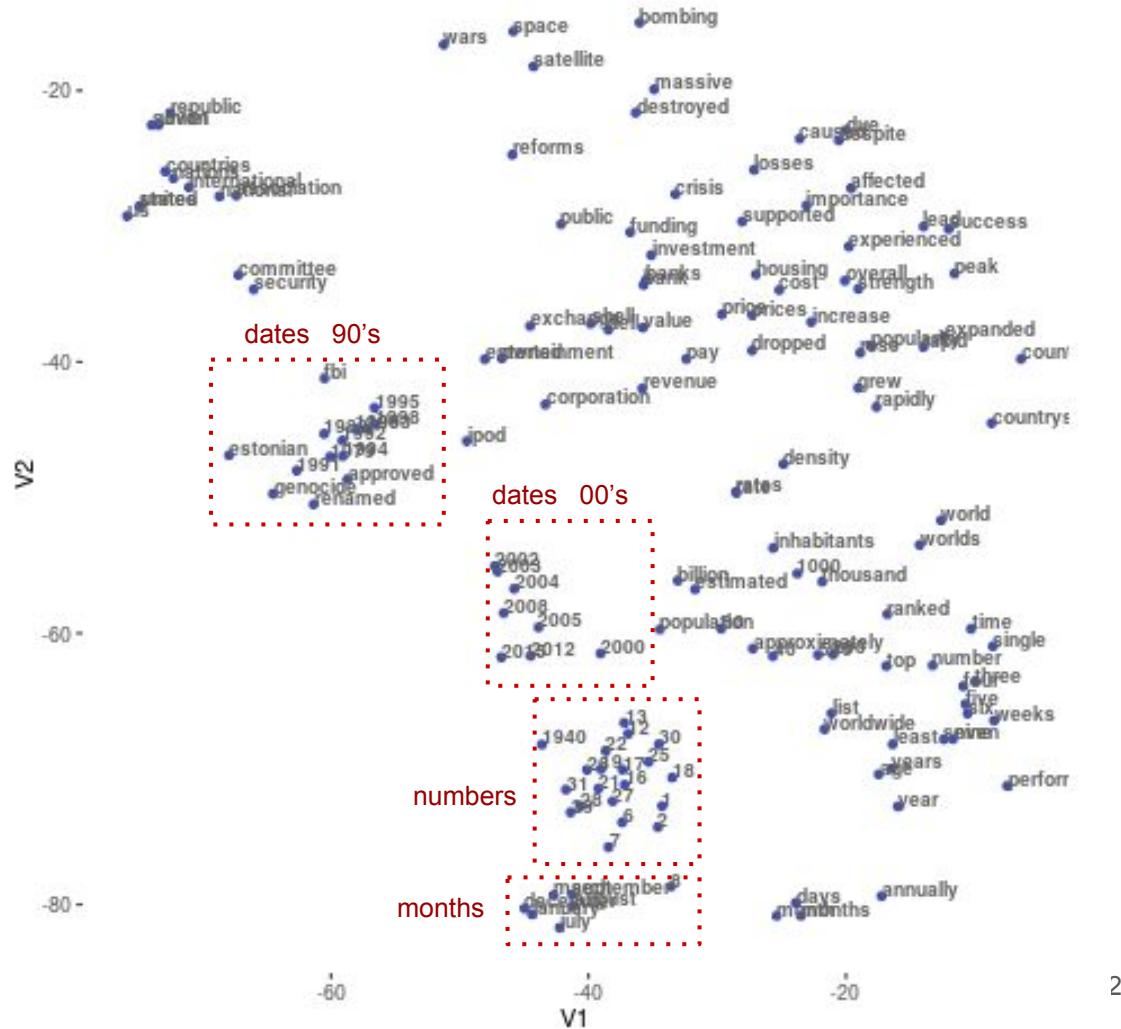


GLOVE Linear Embedding

Window Size = 15

Vector Size = 100

Cluster = 2

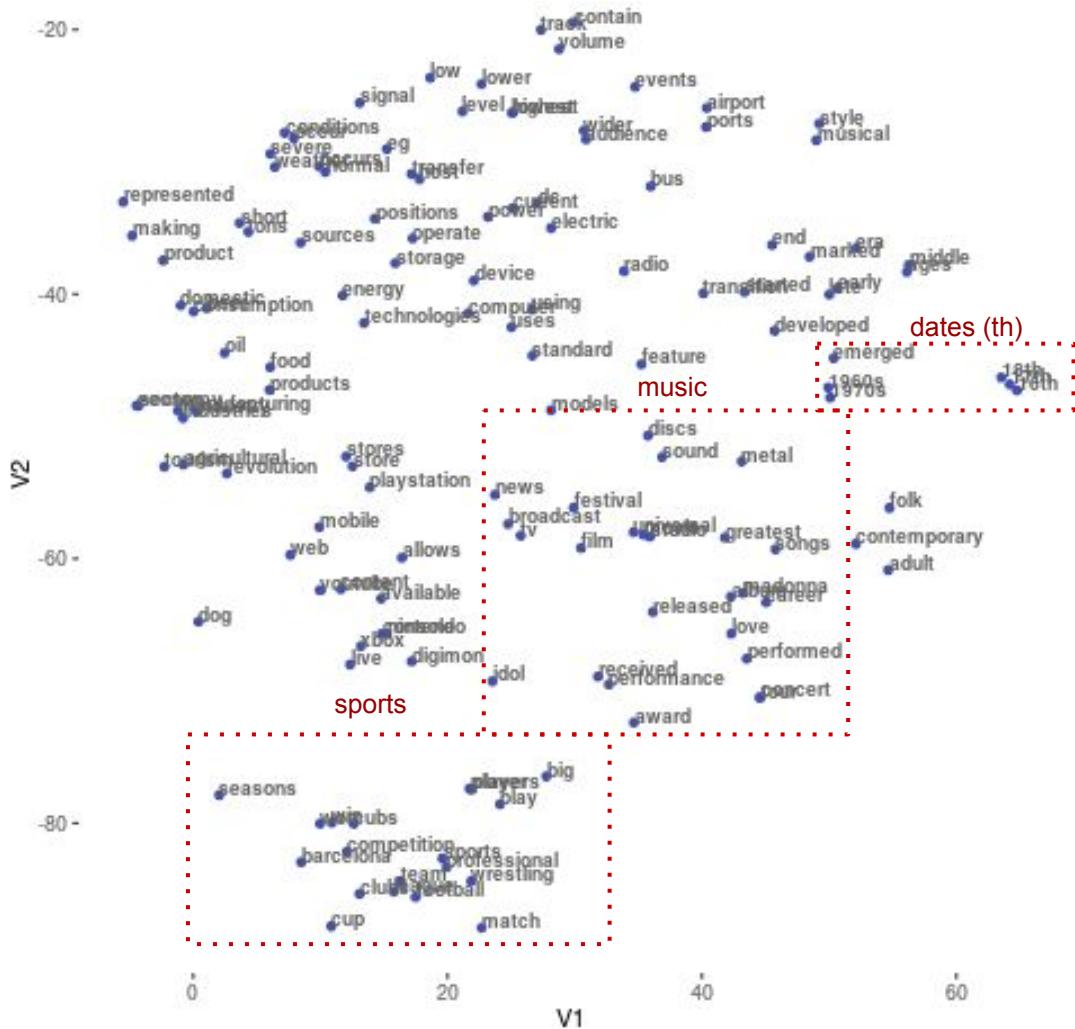


GLOVE Linear Embedding

Window Size = 15

Vector Size = 100

Cluster = 4

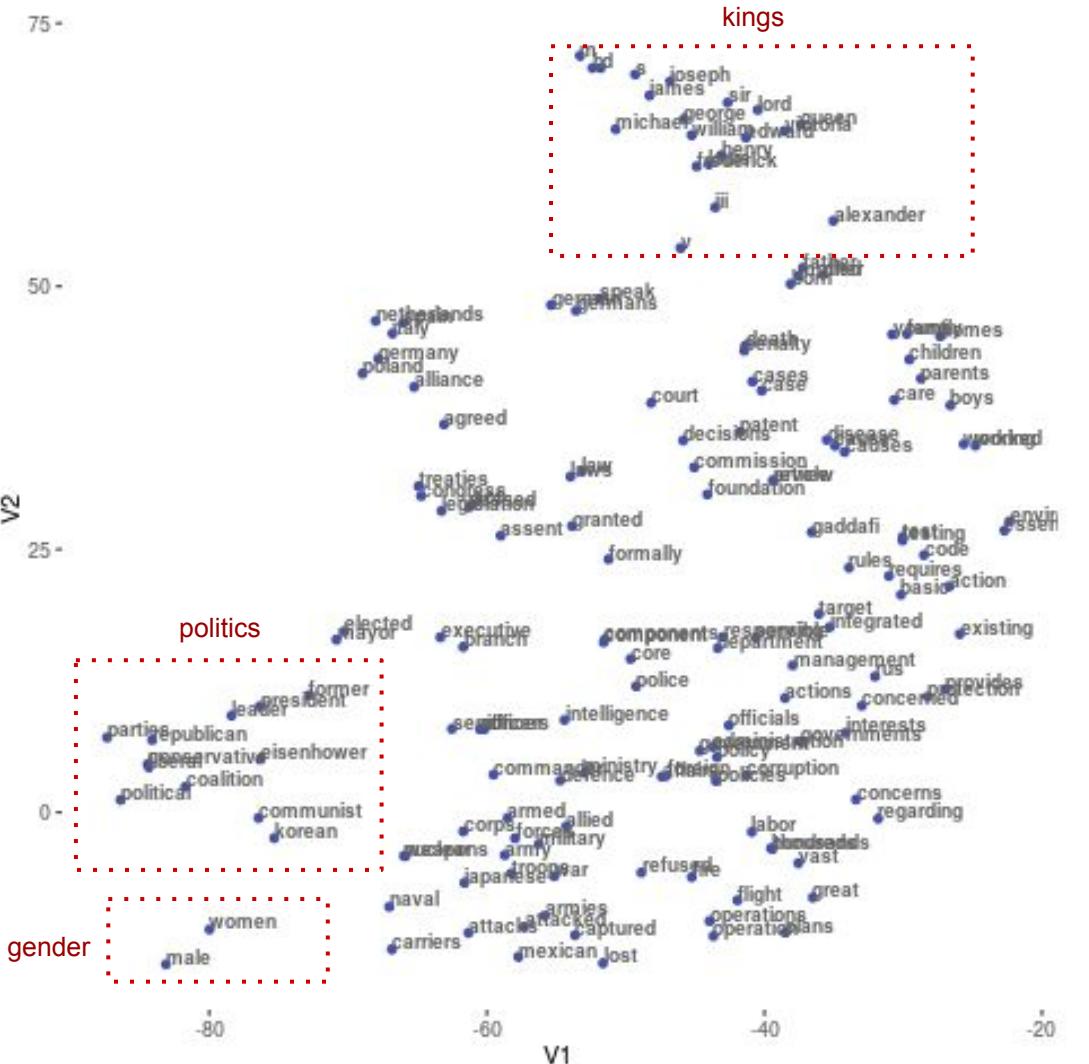


GLOVE Linear Embedding

Window Size = 15

Vector Size = 100

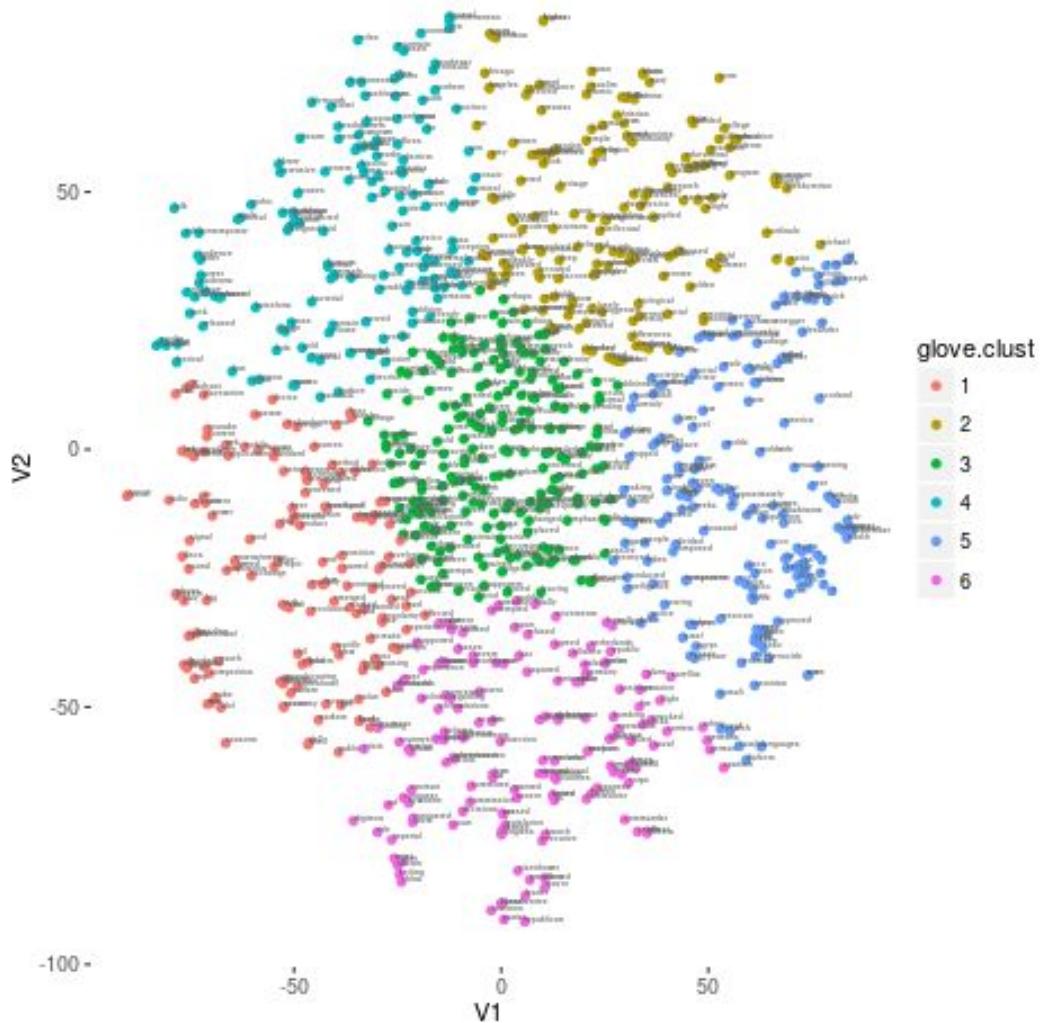
Cluster = 5



GLOVE Linear Embedding

Window Size = 20

Vector Size = 100

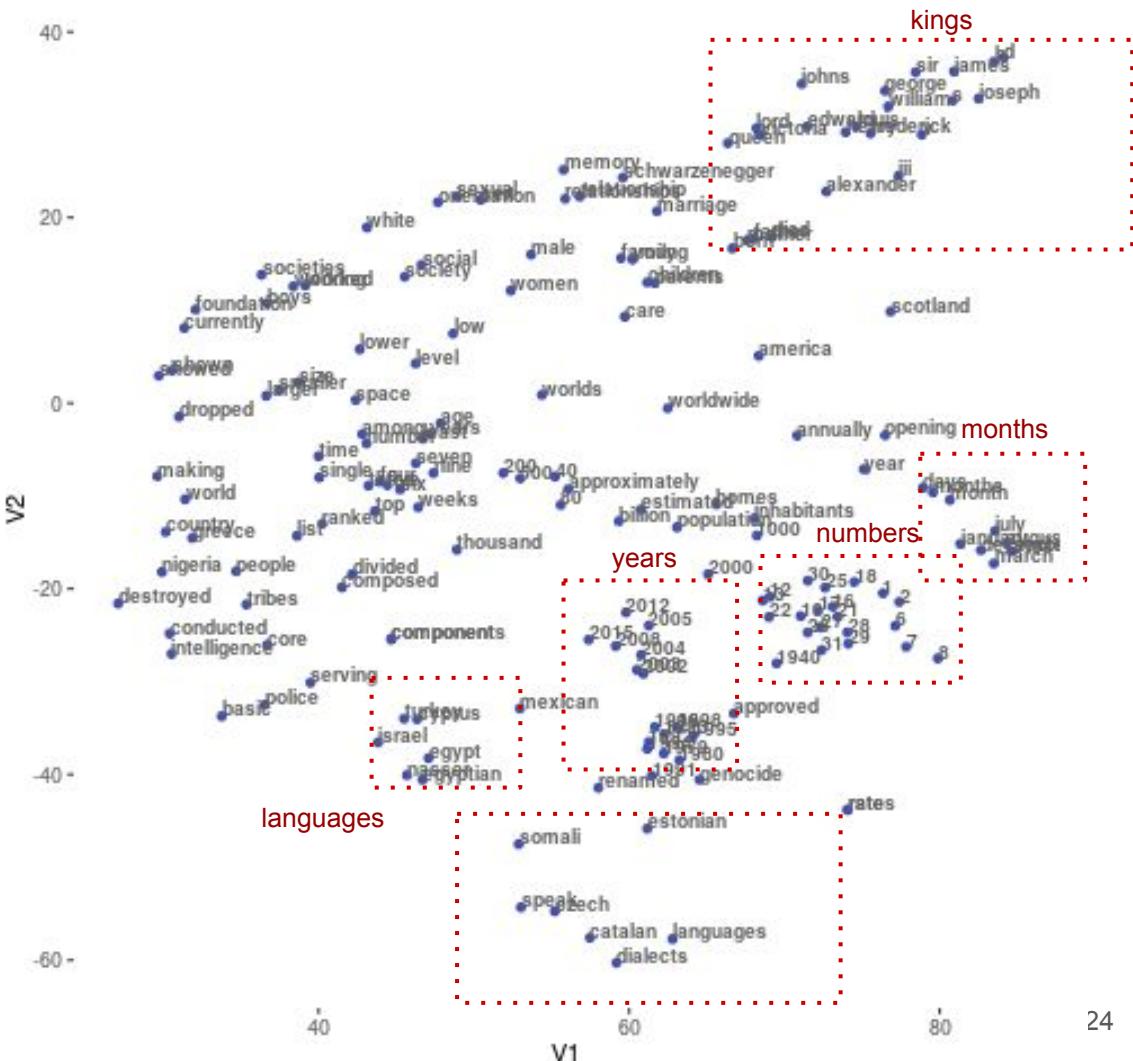


GLOVE Linear Embedding

Window Size = 20

Vector Size = 100

Cluster = 5

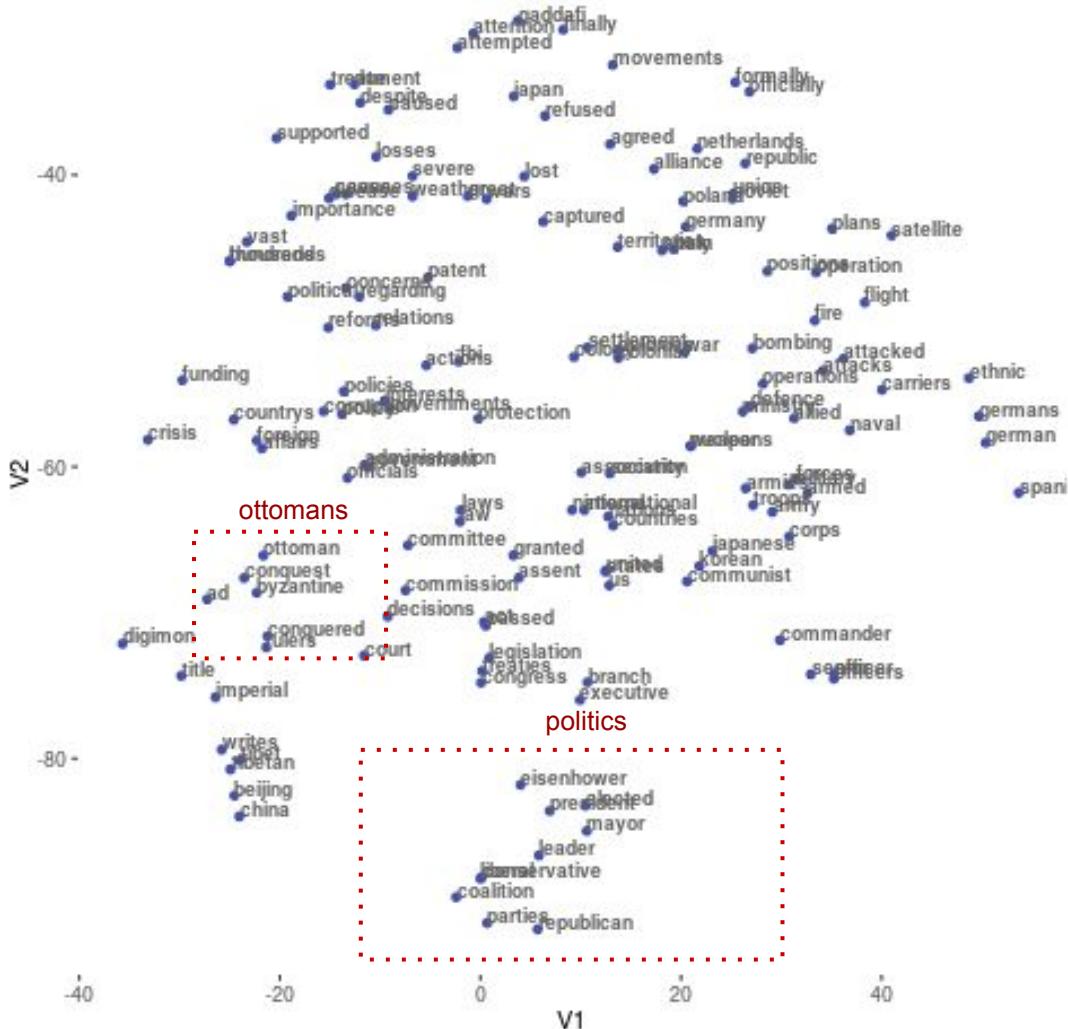


GLOVE Linear Embedding

Window Size = 20

Vector Size = 100

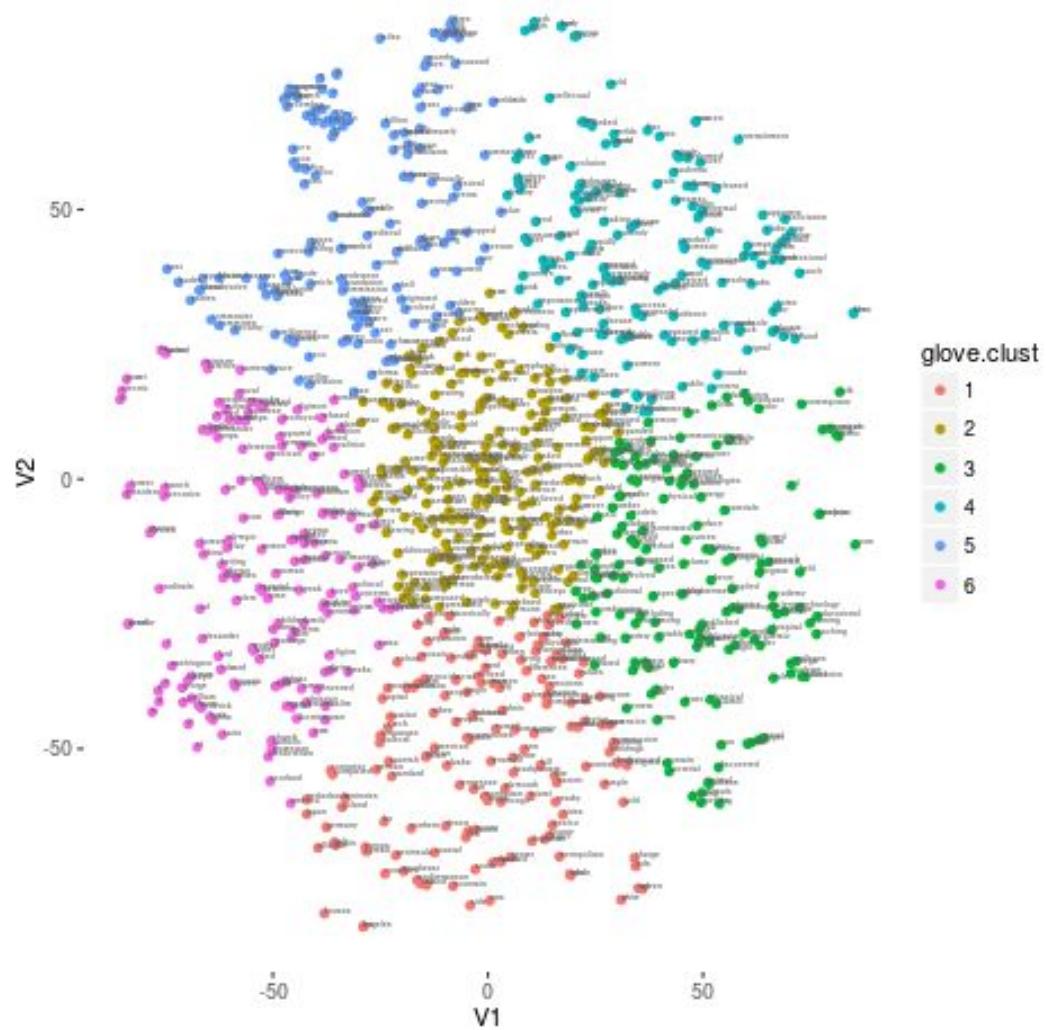
Cluster = 6



GLOVE Linear Embedding

Window Size = 5

Vector Size = 100

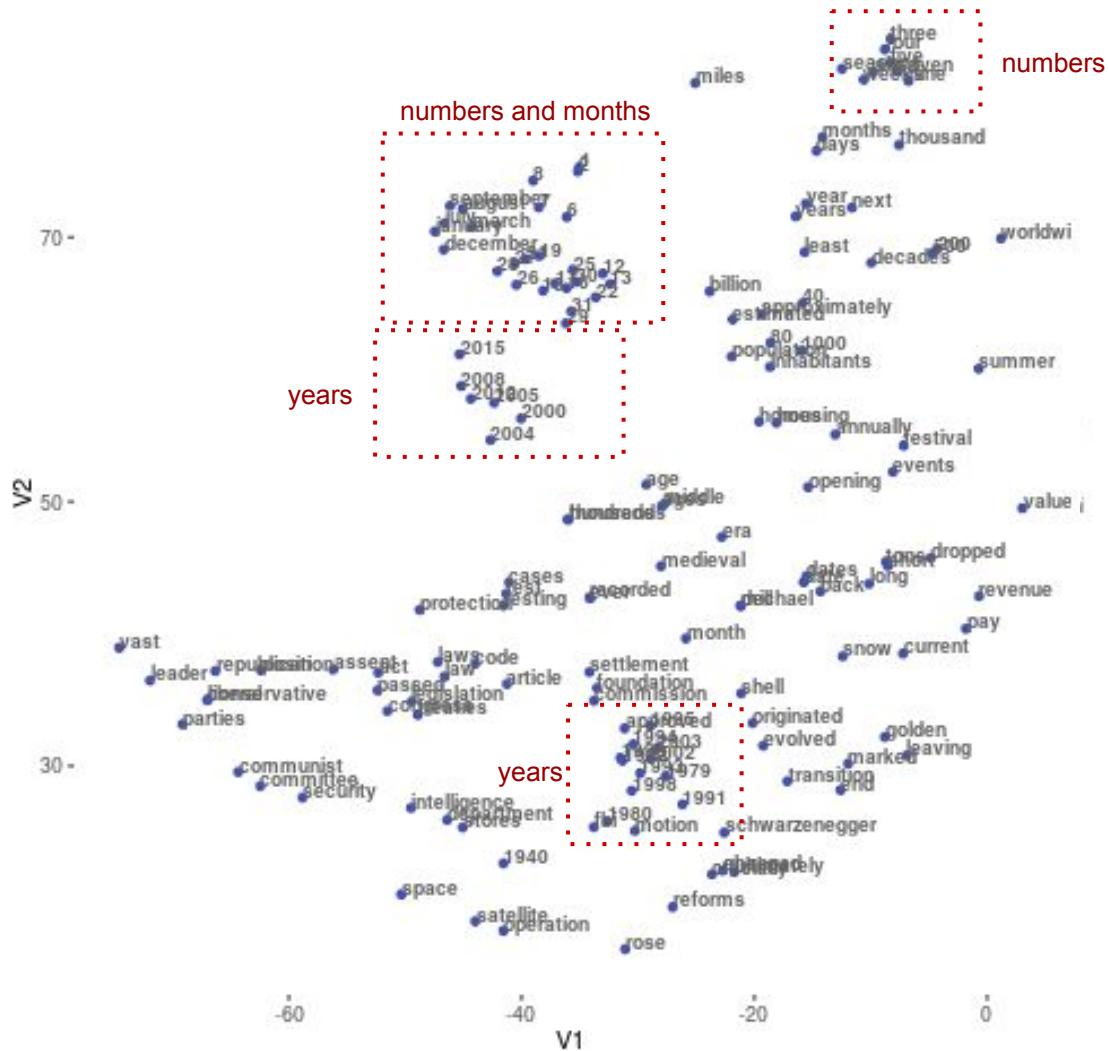


GLOVE Linear Embedding

Window Size = 5

Vector Size = 100

Cluster = 5

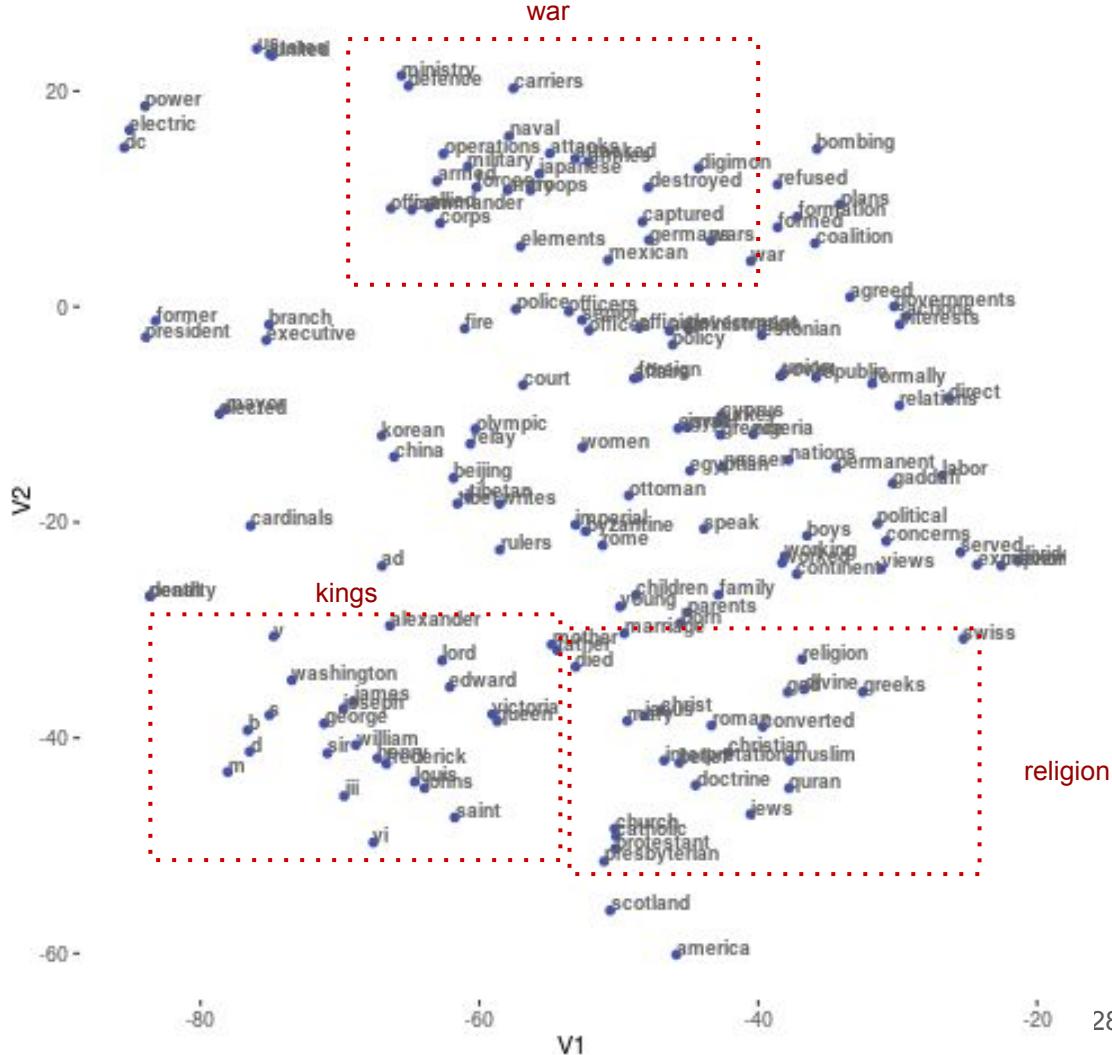


GLOVE Linear Embedding

Window Size = 5

Vector Size = 100

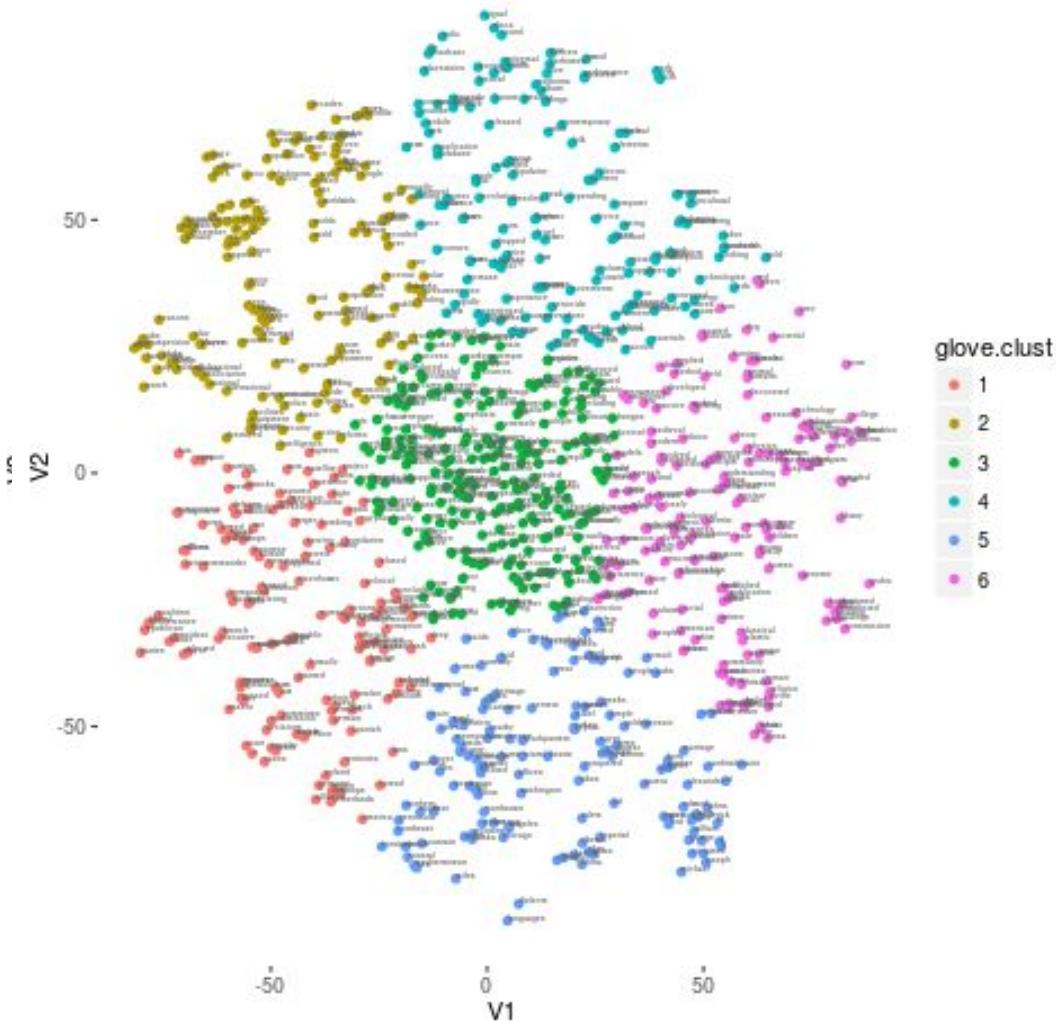
Cluster = 6



GLOVE Linear Embedding

Window Size = 15

Vector Size = 500

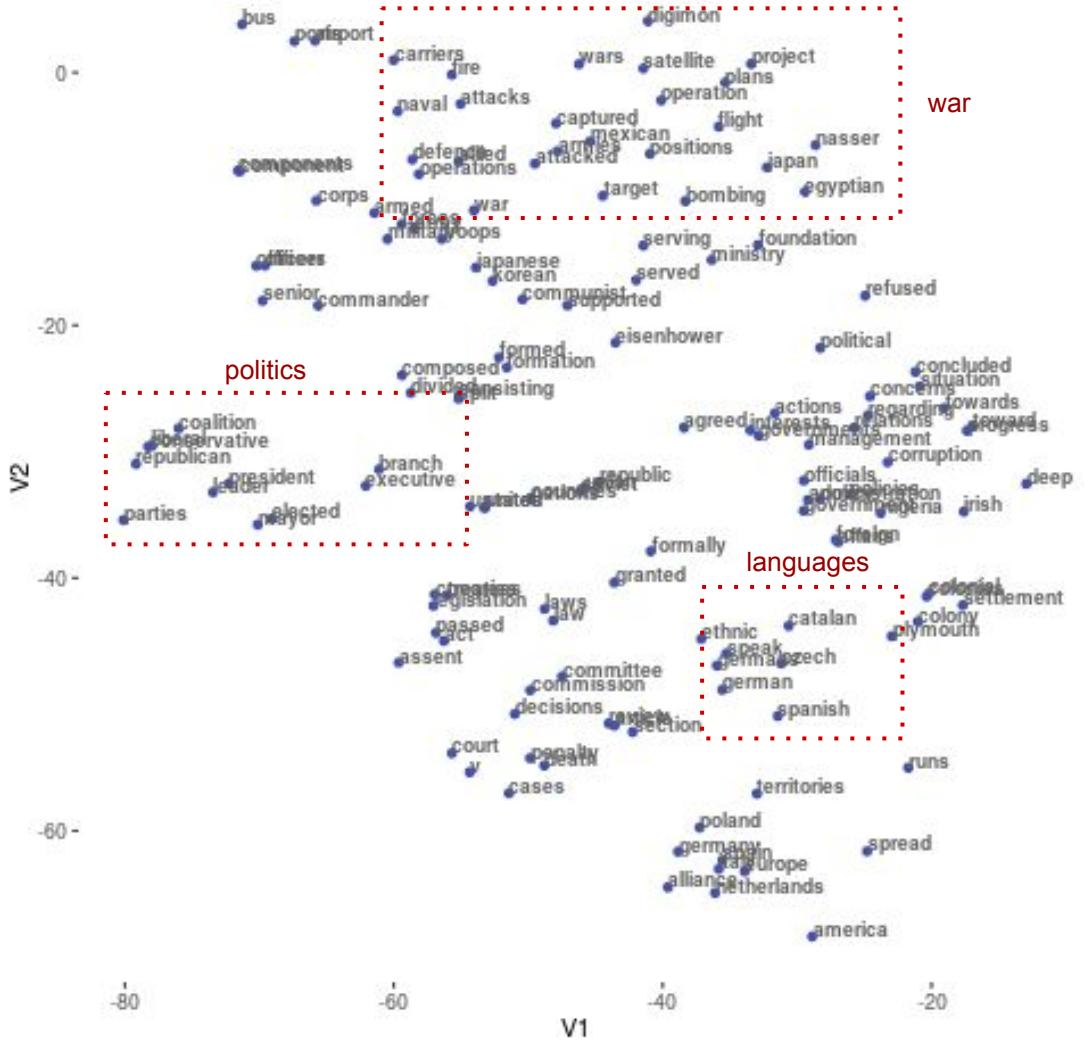


GLOVE Linear Embedding

Window Size = 15

Vector Size = 500

Cluster = 1

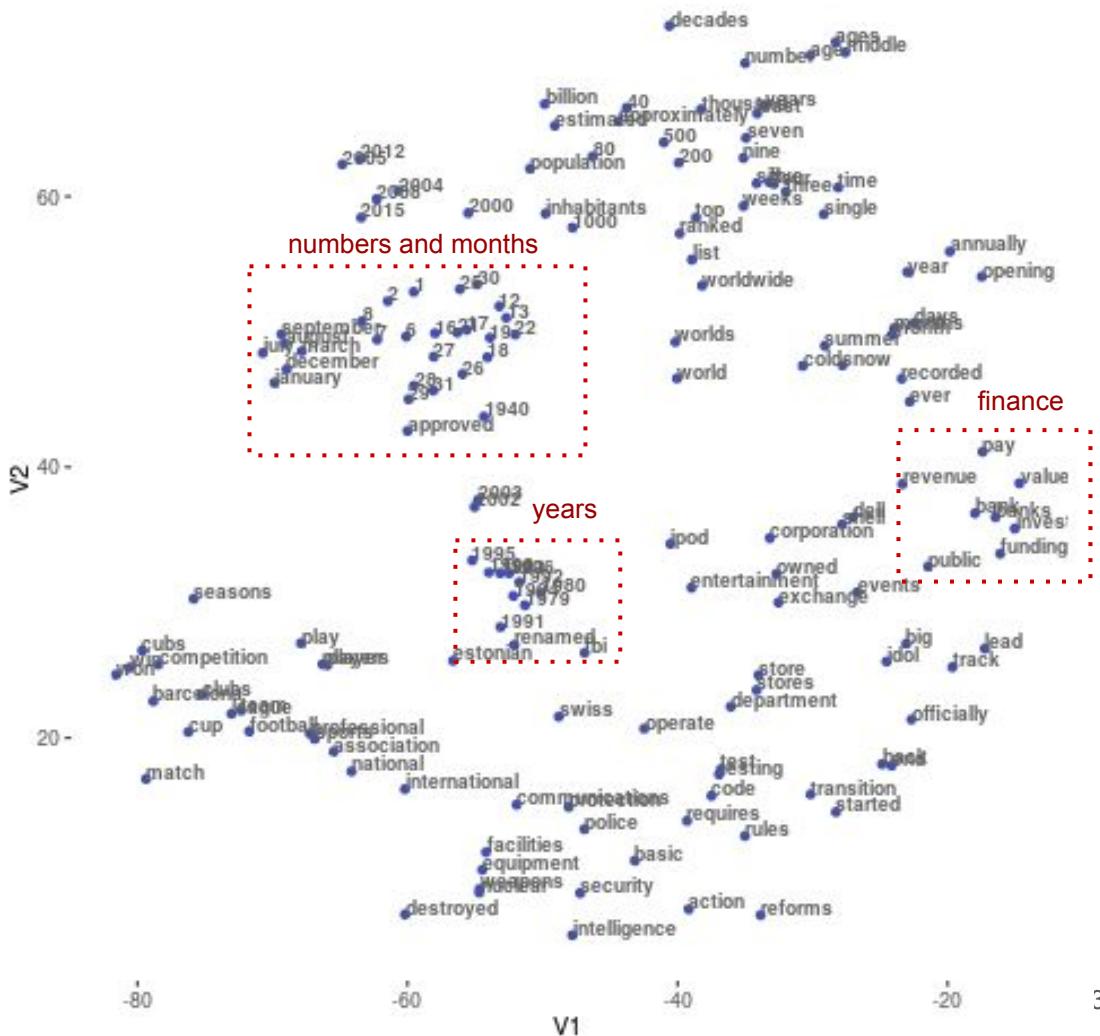


GLOVE Linear Embedding

Window Size = 15

Vector Size = 500

Cluster = 2



GLOVE

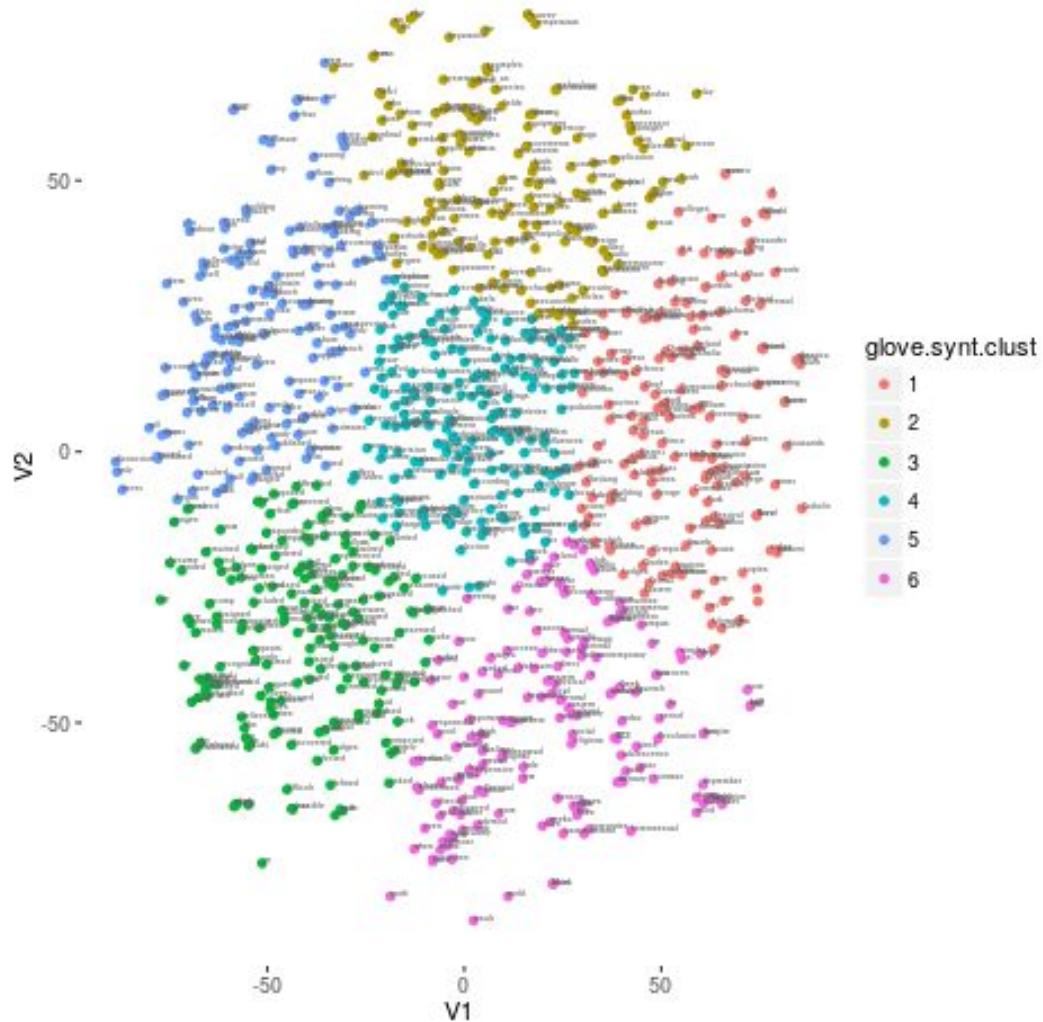
Syntactic Embedding

- Min words in voc = 100
- Size of vectors = 100, 300, 500
- Size of window = 5, 15, 20

GLOVE Syntactic Embedding

Window Size = 15

Vector Size = 300



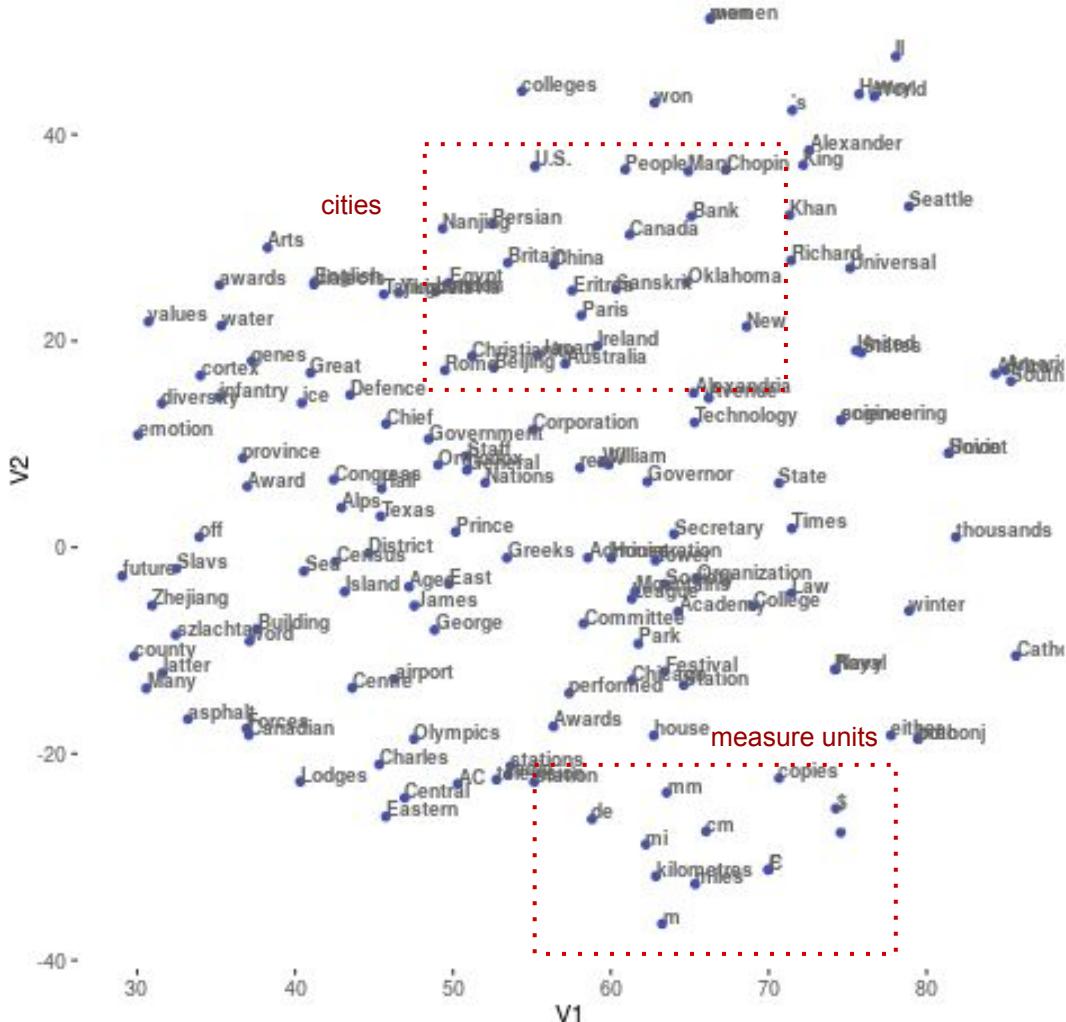
GLOVE Syntactic Embedding

Window Size = 15

Vector Size = 300

Cluster = 1

Captures different relations



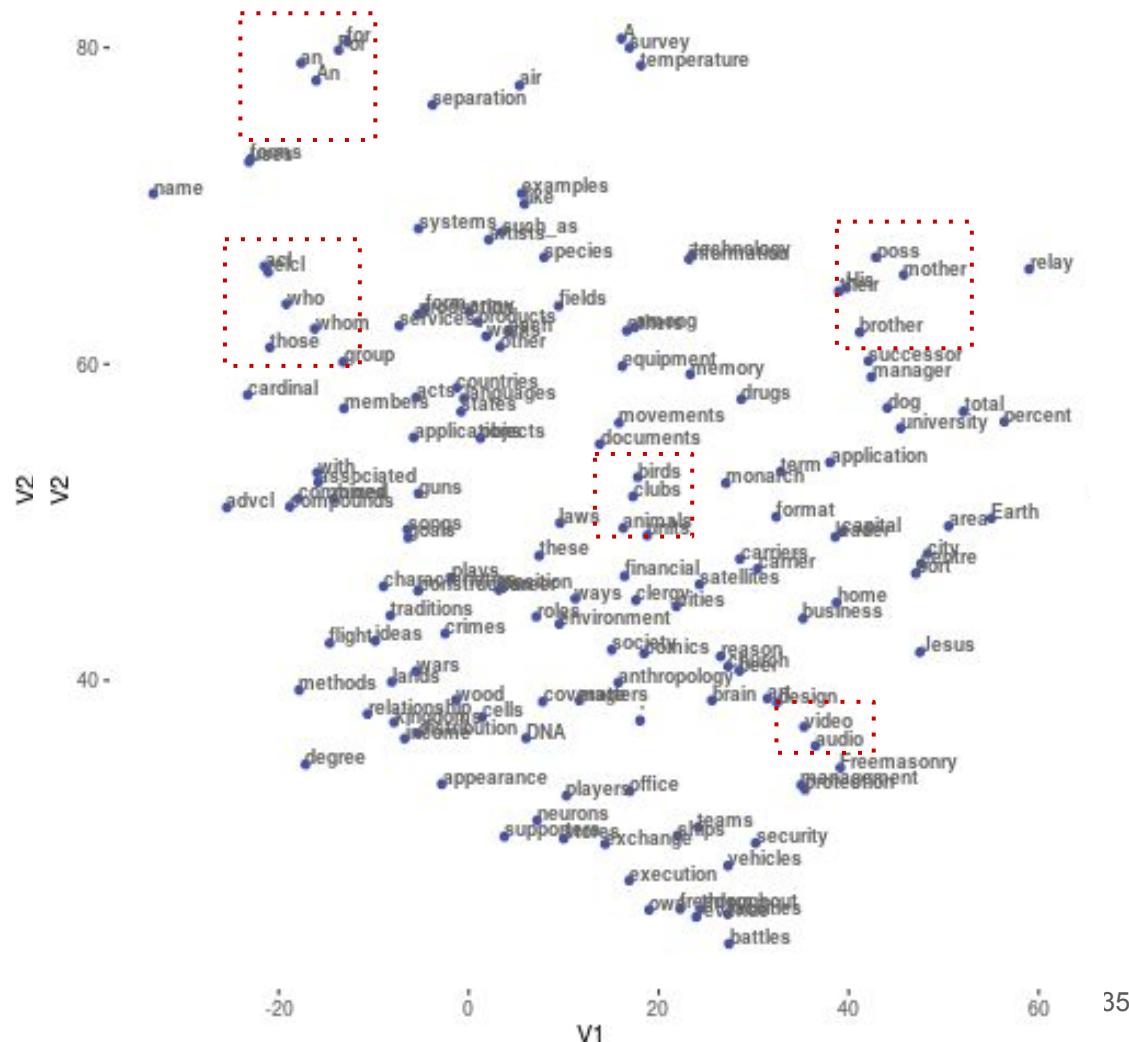
GLOVE Syntactic Embedding

Window Size = 15

Vector Size = 300

Cluster = 2

Captures different relations



GLOVE Topic Embedding

Topic 1th:

- jewish 0.022814
- jews 0.021276
- communities 0.009680
- see 0.005708
- judaism 0.005644
- orthodox 0.005516
- community 0.005324
- hebrew 0.005068
- israel 0.003658
- palestine 0.001864
- synagogue 0.001544
- persecution 0.001416
- jerusalem 0.001352
- group 0.001224
- holocaust 0.001224
- judah 0.001160

Topic 5th:

- pope 0.014170
- paul 0.008777
- john 0.006652
- cardinal 0.006597
- cardinals 0.005726
- bishops 0.005508
- athanasius 0.005344
- vi 0.005072
- rome 0.004963
- bishop 0.004309
- pius 0.003819
- see 0.003547
- vatican 0.003492
- papal 0.003056
- order 0.003002
- saint 0.002675

Topic 9th:

- economic 0.013044
- financial 0.009602
- economy 0.008634
- government 0.008365
- development 0.008311
- industry 0.007559
- public 0.007317
- world 0.005945
- trade 0.005918
- also 0.005649
- international 0.005596
- countries 0.005569
- production 0.005112
- sector 0.004762
- crisis 0.004762
- organization 0.004708

GLOVE

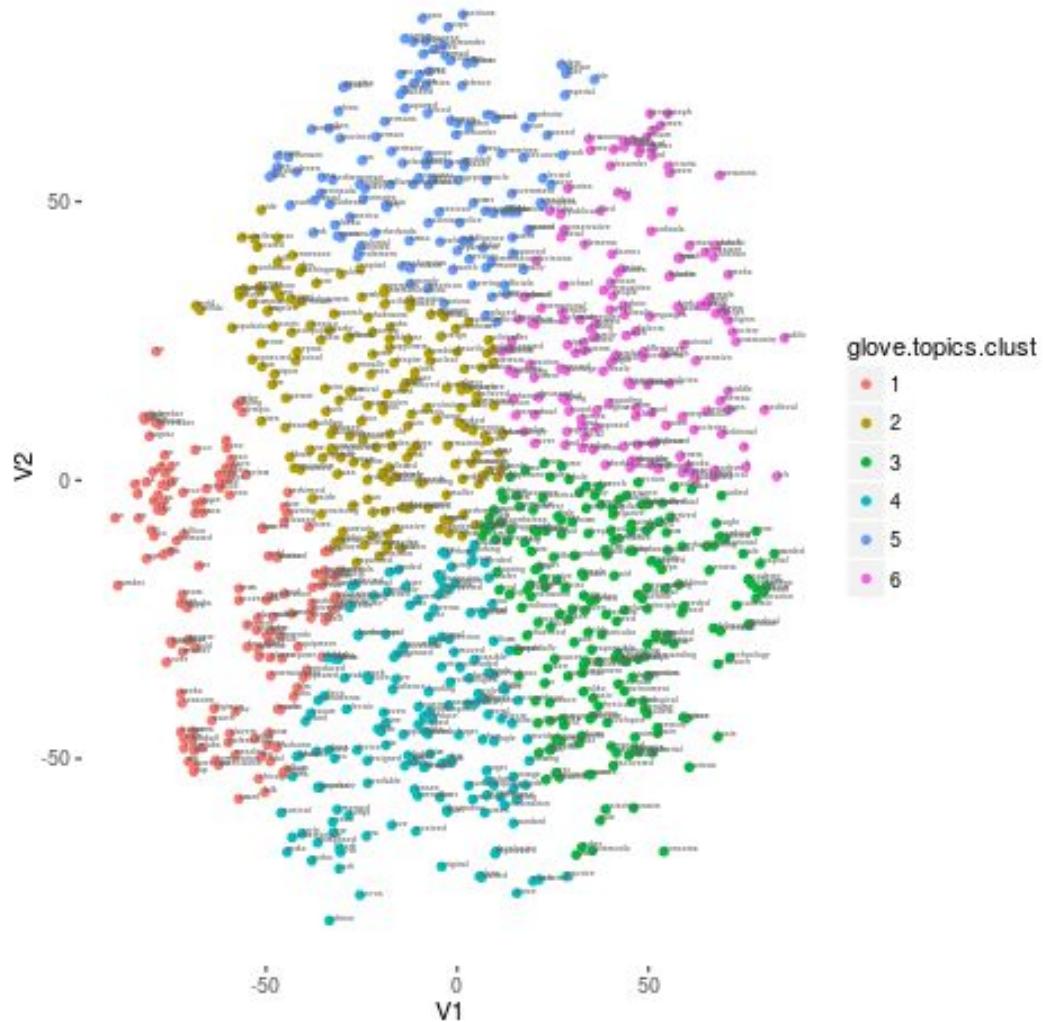
Topic Embedding

- Min words in voc = 100
- Size of vectors = 300
- Size of window = 15

GLOVE Topic Embedding

Window Size = 15

Vector Size = 300



GLOVE Topic Embedding

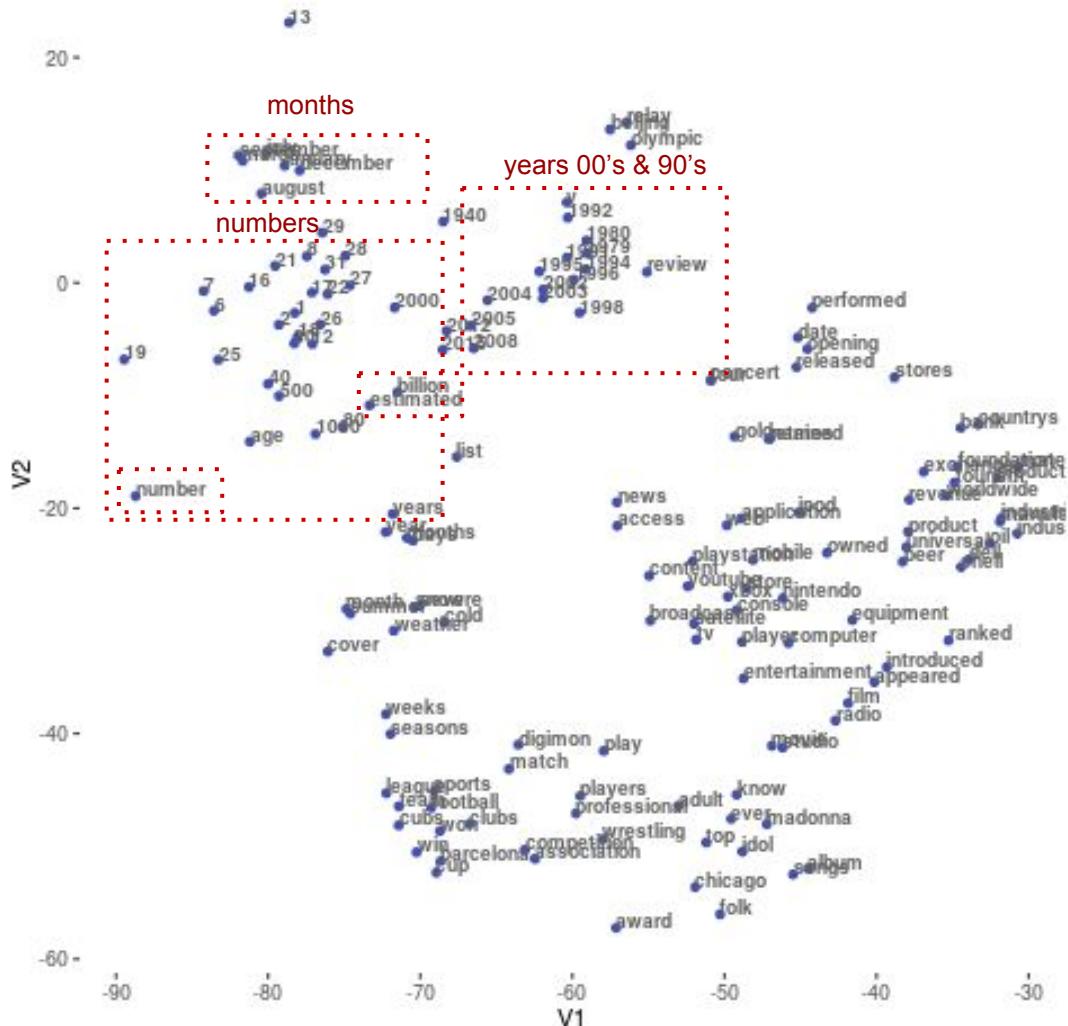
Window Size = 15

Vector Size = 300

Cluster = 1

GLOVE with topic capture

Broad topics



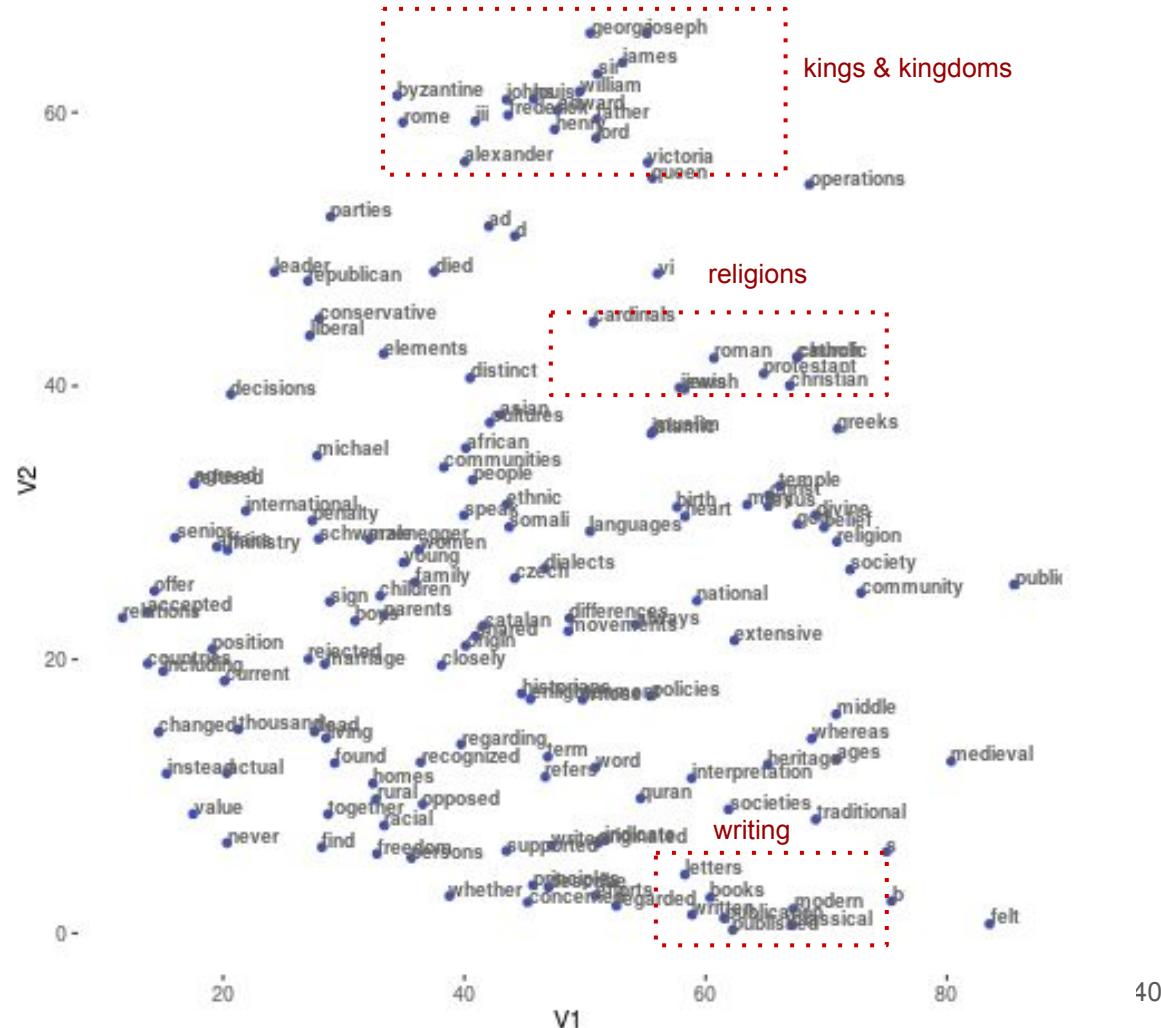
GLOVE Topic Embedding

Window Size = 15

Vector Size = 300

Cluster = 6

GLOVE with topic capture
Broad topics



Exploratory analysis

- General statistics
- Lexical analysis
- Syntactic analysis
 - Embeddings
 - Word
 - Sentence
 - Paragraph

Similar words to 6W+how questions

Python's Doc2Vec on questions

Min_count=10

Window Size = 10

Vector Size = 100

what (80% of questions)

1. which: 0.67
2. where
3. represent
4. resemble
5. supports
6. origins
7. institution
8. protect
9. formal
10. mainly

who

1. succeeded: 0.78
2. successor
3. supports
4. prevented
5. group
6. party
7. freemasons
8. criticized
9. rebel
10. toward

Similar words to 6W+how questions

how

1. there: 0.73
2. about
3. people
4. lines
5. live
6. days
7. million
8. many
9. millions
10. killed

which

1. named: 0.67
2. dominated
3. consisted
4. formed
5. mayor
6. divides
7. Somali
8. dominant
9. formerly
10. reform

Similar words to 6W+how questions

when

1. why: 0.71
2. son
3. John
4. succeeded
5. leave
6. revolution
7. richard
8. constantinople
9. ask
10. before

where

1. v

Similar words to 6W+how questions

why

1. stepper: 0.82
2. absorb
3. doing
4. mark
5. without
6. efficacy
7. genes
8. can
9. insects
10. maintain

Exploratory analysis

- General statistics
- Lexical analysis
- Syntactic analysis
 - Embeddings
 - Word
 - Sentence
 - Paragraph

Paragraph embeddings detect similarities between words

Python's Doc2Vec on paragraphs

Min_count=10

Window Size = 10

Vector Size = 100

Synonym identification:

- $\text{sim}([\text{'college'}, \text{'professor'}], [\text{'university'}, \text{'teacher'}]) = \text{0.92}$
- $\text{sim}([\text{'marriage'}, \text{'husband'}, \text{'baby'}], [\text{'wife'}, \text{'wedding'}, \text{'children'}]) = \text{0.85}$
- $\text{sim}([\text{'house'}, \text{'residence'}, \text{'bed'}, \text{'accommodation'}, \text{'address'}], [\text{'shelter'}, \text{'mansion'}, \text{'home'}, \text{'place'}]) = \text{0.77}$

This analysis also detects non-related terms and analogies

Python's Doc2Vec on paragraphs

Min_count=10

Window Size = 10

Vector Size = 100

Non-related terms identification:

- similarity('husband', 'floor') = **0.30**
- similarity('night', 'chicken') = **0.29**
- similarity('computer', 'city') = **0.22**

Analogies

- woman is to king as man is to...? prince
- Most similar to “queen”: Madonna, widow, performed
- Most similar to “man”: said, wrote, god

The topics found with LDA can be refined using paragraph embeddings

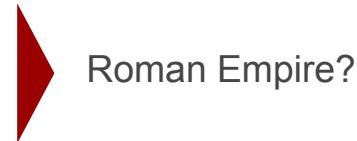
LDA:

church + roman + first + emperor ~ **history**

state + govern + force + war ~ **government**

Most similar words to LDA keywords:

1. rome: 0.86
2. byzantine
3. centuries
4. patriarch
5. 14th
6. survived
7. 12th
8. successors
9. constantine
10. succession



1. government: 0.85
2. administration
3. sovereign
4. military
5. suppress
6. forces
7. initiated
8. supported
9. organized
10. urged



The topics found with LDA can be refined using paragraph embeddings

LDA:

city + new + state + area + unit ~ **nation-state** game + team + play ~ **sports**

Most similar words to LDA keywords:

1. located: 0.86
2. metropolitan
3. headquarters
4. county
5. designated
6. operated
7. downtown
8. currently
9. main
10. serves

metropolitan areas?

1. championship: 0.89
2. games
3. players
4. fans
5. exhibition
6. afl
7. matches
8. teams
9. nfl
10. super

championship?

The topics found with LDA can be refined using paragraph embeddings

LDA:

music + film + record ~ art

Most similar words to LDA keywords:

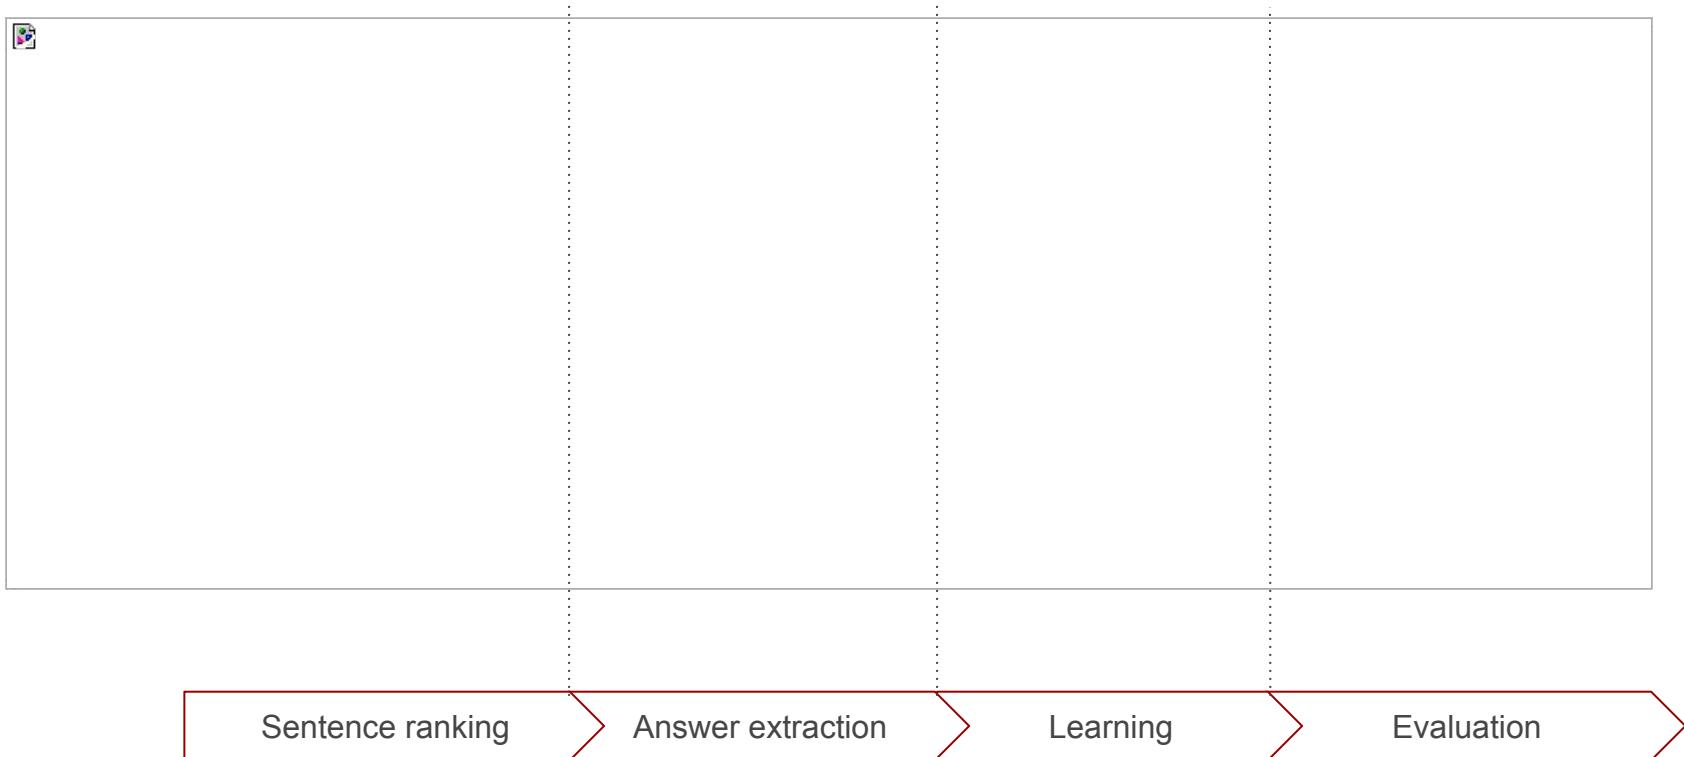
1. films: 0.9
2. featured
3. movie
4. studio
5. singers
6. guitar
7. songs
8. artist
9. albums
10. hip-hop



music and film
recording?

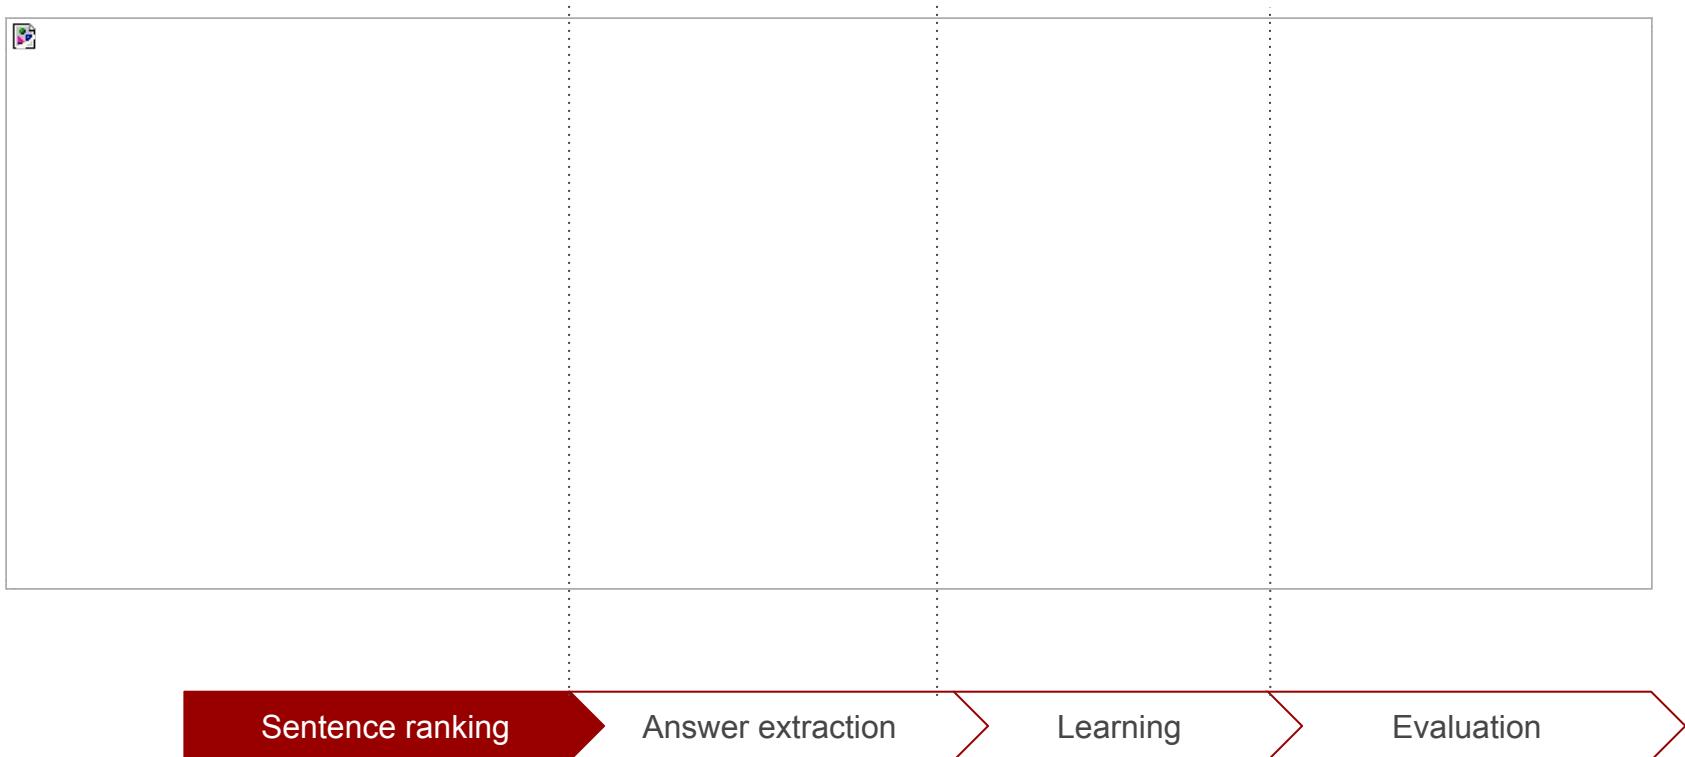
Pipeline description

High level baseline pipeline



Sentence ranking

High level baseline pipeline



Sentence Ranking

The whole idea of sentence ranking is to exploit lexical and syntactical similarities between the question and the answer passage to obtain the sentence with the highest likelihood of being the answer.

Sentence Ranking

Convolutional Neural Networks

Convolutional neural network model for
reranking pairs of short texts:

- Learn optimal vector representation of Q-D
- Learn a similarity function between Q-D vectors

Sentence Ranking

Convolutional Neural Networks

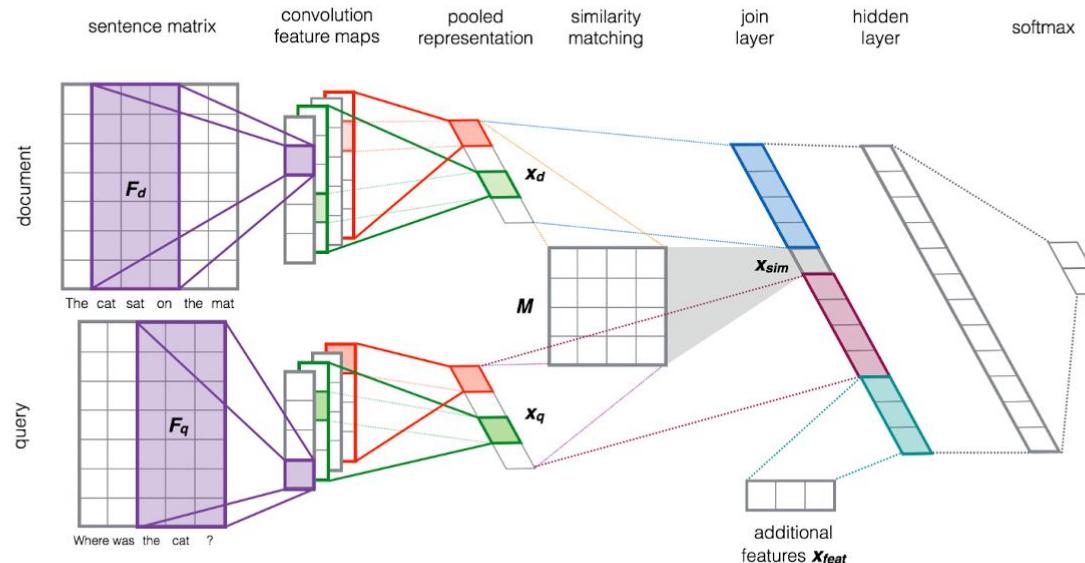


Figure 2: Our deep learning architecture for reranking short text pairs.

Sentence Ranking

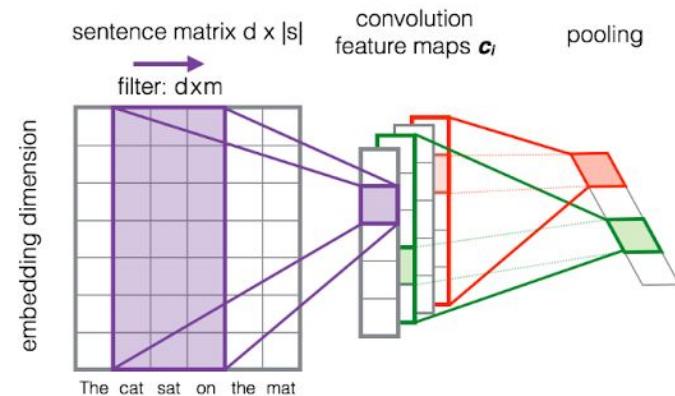
Convolutional Neural Networks

Sentences are represented as sequences of words, where each word is an $|s|$ dimensional continuous representation.

$$\mathbf{S} = \begin{bmatrix} | & | & | \\ \mathbf{w}_1 & \dots & \mathbf{w}_{|s|} \\ | & | & | \end{bmatrix}$$

A filter \mathbf{f} is applied to the sequence in order to capture interactions among words.

$$\mathbf{c}_i = (\mathbf{s} * \mathbf{f})_i = \mathbf{s}_{[i-m+1:i]}^T \cdot \mathbf{f} = \sum_{k=i}^{i+m-1} s_k f_k$$

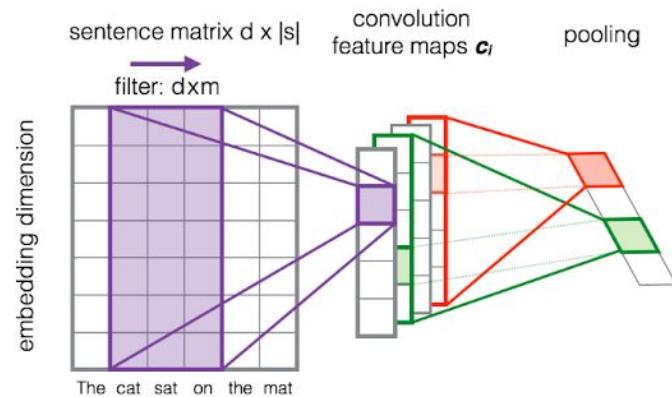


Sentence Ranking

Convolutional Neural Networks

After this is done, a nonlinear activation function, ReLU in this case, is applied to every c_i and the results are pooled together via max pooling into a single c_{pooled} array representation.

$$\mathbf{c}_{\text{pooled}} = \begin{bmatrix} \text{pool}(\alpha(\mathbf{c}_1 + b_1 * \mathbf{e})) \\ \dots \\ \text{pool}(\alpha(\mathbf{c}_n + b_n * \mathbf{e})) \end{bmatrix}$$

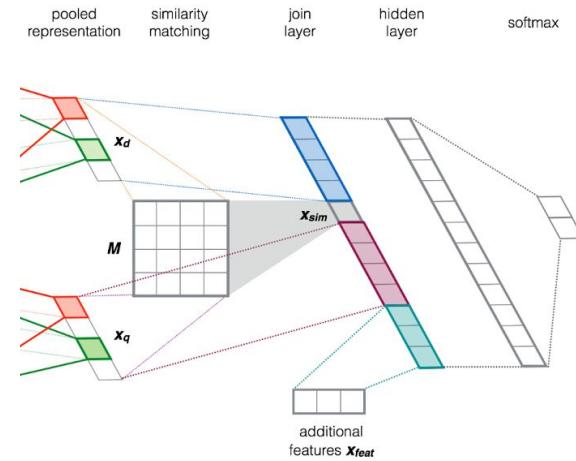


Sentence Ranking

Convolutional Neural Networks

Once these representations are obtained for each sentence x_d and each query x_q , a x_{sim} score is obtained by $x_d' M x_q$ and an x_{join} is created by simple concatenation.

Each x_{join} is passed through a hidden layer to exploit interactions among its different components, and finally a softmax is used for the ranking.



Sentence Ranking

Convolutional Neural Networks

The model is trained to minimize the cross-entropy function:

$$\begin{aligned}\mathcal{C} &= -\log \prod_{i=1}^N p(y_i | \mathbf{q}_i, \mathbf{d}_i) + \lambda \|\theta\|_2^2 \\ &= -\sum_{i=1}^N [y_i \log \mathbf{a}_i + (1 - y_i) \log(1 - \mathbf{a}_i)] + \lambda \|\theta\|_2^d,\end{aligned}$$

where \mathbf{a} is the output of the softmax and θ contains all the parameters of the network:

$$\theta = \{\mathbf{W}; \mathbf{F}_q; \mathbf{b}_q; \mathbf{F}_d; \mathbf{b}_d; \mathbf{M}; \mathbf{w}_h; b_h; \mathbf{w}_s; b_s\}$$

Regularization is used to avoid overfitting and stochastic gradient descent to cope with the non convexity of the problem.

Sentence Ranking

Convolutional Neural Networks

In our model, we added an hybrid vector representation that used both, the representation trained over the AQUAINT corpus (to obtain the most general context of each word), and over the SQuAD dataset (to obtain the particular uses of each word). We also used Jaccard similarity as a proxy of relevance judgments, and we added topic information to the x_{joint} representation.

Sentence Ranking

BM25 & Jaccard similarity

Another approach that uses only lexical similarity, under the bag of words was applied, namely BM25 and Jaccard similarity:

- BM25

$$\text{score}(D, Q) = \sum_{i=1}^n \text{IDF}(q_i) \cdot \frac{f(q_i, D) \cdot (k_1 + 1)}{f(q_i, D) + k_1 \cdot \left(1 - b + b \cdot \frac{|D|}{\text{avgdl}}\right)}$$
$$\text{IDF}(q_i) = \log \frac{N - n(q_i) + 0.5}{n(q_i) + 0.5},$$

- Jaccard similarity

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|}$$

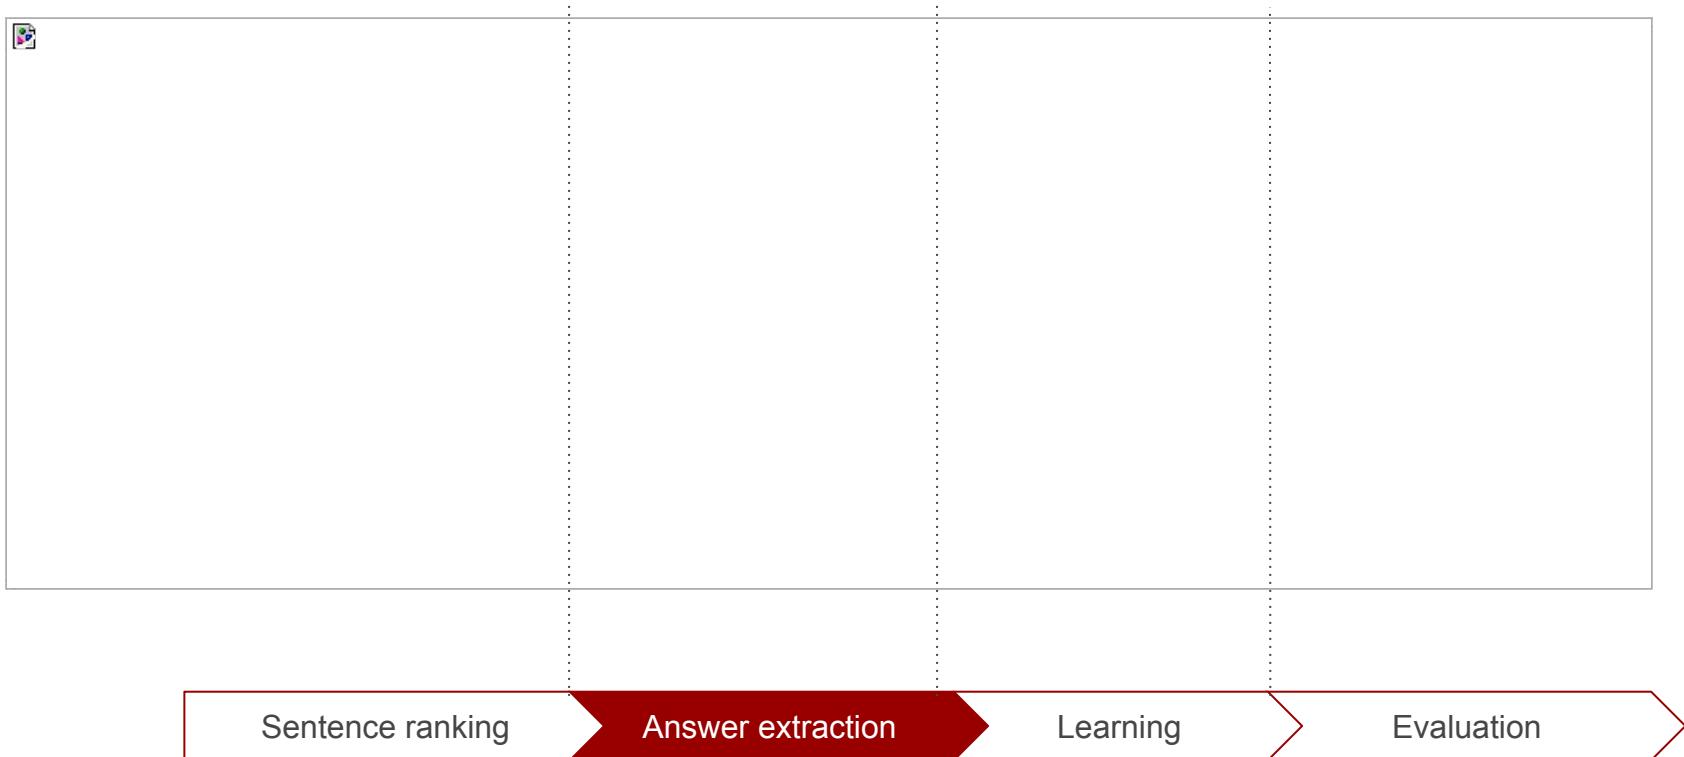
Sentence Ranking

Results

MODEL	MRR Score
Convolutional Neural Networks	.25
BM25	.71
Jaccard	.76

We believe that the bad results of the ConvNets is due to underfitting of the training data.

High level baseline pipeline



Answer extraction

Answer extraction

Random Forests

Idea: Use features that extract lexical, syntactical and semantical structure of sentence, question and answer to train a classifier.

For each word in candidate answer sentence:

- indicator of right neighbor in question
- right neighbor NER
- right neighbor POS
- word Animacy
- word Gender
- word NER
- word Number
- word POS
- word type
- dependency with father
- indicator father in question
- father NER
- father POS
- indicator of word in question
- indicator of left neighbor in question
- left neighbor NER
- left neighbor POS
- question type

Answer extraction

Random Forests

Example: **It**

```
(False, u'It', u'PRP', u'O', False, 'whom', "", "", "", "", u'is', u'VBZ', u'O', False, False, u'replica',  
u'NN', u'O', u'nsubj', False, False, u'INANIMATE', u'SINGULAR', u'NEUTRAL',  
u'PRONOMINAL')
```

Answer extraction Random Forests

Model

- 1000 trees
- 5 variables per cut
- Gini criterion

Results

Training:

F1 score = .49

precision = .62

Test:

F1 score = .47

precision = .6

Answer extraction Random Forests

```
"56d601e41c85041400946ed0": "sacked him seven times and",
"56d601e41c85041400946ed1": "Bowl 50 and",
"56d601e41c85041400946ed2": "tackles 21/2 sacks",
"56d98b33dc89441400fdb53b": "him seven times",
"56d98b33dc89441400fdb53c": "Bowl 50 and",
"56d98b33dc89441400fdb53d": "two",
"56d98b33dc89441400fdb53e": "tackles 21/2 sacks",
"56be5333acb8001400a5030a": "Bowl 50 in",
"56be5333acb8001400a5030b": "of 5 million",
"56be5333acb8001400a5030c": "and Bruno Mars who headlined",
"56be5333acb8001400a5030d": "Bruno Mars who",
"56be5333acb8001400a5030e": "Bruno Mars who",
"56beaf5e3aeaaa14008c91fd": "50",
"56beaf5e3aeaaa14008c91fe": "average of 5 million for",
"56beaf5e3aeaaa14008c91ff": "Bruno Mars who",
"56beaf5e3aeaaa14008c9200": "Mars",
"56beaf5e3aeaaa14008c9201": "thirdmost",
"56bf1ae93aeaaa14008c951b": "Bowl 50 in",
"56bf1ae93aeaaa14008c951c": "of 5 million",
"56bf1ae93aeaaa14008c951e": "Bruno Mars who",
"56bf1ae93aeaaa14008c951f": "Bruno Mars who",
"56d2051ce7d4791d00902608": "of 5 million",
"56d2051ce7d4791d00902609": "of 5 million",
"56d2051ce7d4791d0090260a": "Bruno Mars who",
"56d2051ce7d4791d0090260b": "and Bruno Mars who headlined",
"56d602631c85041400946ed8": "50",
"56d602631c85041400946eda": "Bruno Mars who"
```

Pipeline implementation

Pipeline implementation

Our pipeline implementation supports:

- An end to end pipelined execution.
- Model training
- Model testing
- Interactive Mode

Pipeline implementation

Model training:

The system allows you to choose the number of sentences to be considered as part of the answer as well as the number of instances used on the training phase.

```
2016-09-12 11:20:59 * latitude in ~/Documents/CMU/pipeline
± | master S:2 U:3 ?:4 X| → ./exec_pipe.sh train.json 1 100
Running sentence ranker...
Sentence ranker done. Output -> ./output/train_rank.json
#####
Create features? choose [y]es|[n]o □
```

Pipeline implementation

Model testing:

It also gives you the option to train or test the model. And provides a final evaluation with Stanford's script.

```
#####
Create features? choose [y]es|[n]o n
Do you want to t[r]ain|t[e]st? e
#####
Running answer extractor...
#####
Answer extraction test done. Output -> ./output/predicts.json
Stanford's test format. Output -> ./output/test_stanford.json

Stanford's evaluation:
{"f1": 0.20952380952380956, "exact_match": 0.08490566037735849}

Done!!!
```

Pipeline implementation

Model interactive mode:

Finally, to enable testing of new models, the system also supports interactive mode.

```
2016-09-12 11:27:24 * | latitude in ~/Documents/CMU/pipeline
± | master S:2 U:3 ?:4 X| → ./exec_pipe.sh

No input file supplied. Assuming interactive mode.
Is this correct [y/n] o y
Introduce a context
Subway Sadie is a comedy-drama film that premiered in New York on September 12, 1926. It was adapted from Mildred Cram's 1925 short story "Sadie of the Desert" and directed by Alfred Santell. The silent film focuses on a relationship between New York salesgirl Sadie Hermann (Dorothy Mackaill) and subway guard Herb McCarthy (Jack Mulhall), who meet on a subway and become engaged. After Sadie receives a promotion, she must choose between her new job and marrying Herb. The cast includes Charles Murray, Peggy Shaw, Gaston Glass, and Bernard Randall. The film began production in May 1926 and was distributed by First National Pictures. Arthur Edeson served as cinematographer, shooting scenes in a nightclub and a casino, and at Cleopatra's Needle in Central Park. Many publications wrote positively of the film, praising its acting and Santell's direction. Today, it remains unclear if a print of Subway Sadie has survived. A poster of the film can be seen at the New York Transit Museum.

Introduce a question
Who was the director of Subway Sadie? □
```

Pipeline implementation

Model interactive mode:

Finally, to enable testing of new models, the system also supports interactive mode.

```
#####
# Running feature constructor ...
# /usr/lib/python2.7/dist-packages/cffi/model.py:526: UserWarning: 'point_conversion_form_t' has no values explicitly defined; next version will refuse to guess which integer type it is meant to be (unsigned/signed, int/long)
#   % self._get_c_name())
#
# Feature extractor done. Output -> ./output/interact_features.json
#####
# Running answer extractor...
#
# The Answer is:
#
# [1] "Sadie"
#
# Done!!!
```

Pipeline implementation

End to end execution results evaluated under Stanford's metric:

```
{"f1": 0.20368373764600187, "exact_match": 0.07547169811320754}
```

	Exact Match		F1	
	Dev	Test	Dev	Test
Random Guess	1.1%	1.3%	4.1%	4.3%
Sliding Window	13.2%	12.5%	20.2%	19.7%
Sliding Win. + Dist.	13.3%	13.0%	20.2%	20.0%
Logistic Regression	40.0%	40.4%	51.0%	51.0%
Human	80.3%	77.0%	90.5%	86.8%

References

- Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient Estimation of Word Representations in Vector Space. In Proceedings of Workshop at ICLR, 2013.
- Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg Corrado, and Jeffrey Dean. Distributed Representations of Words and Phrases and their Compositionality. In Proceedings of NIPS, 2013.
- Tomas Mikolov, Wen-tau Yih, and Geoffrey Zweig. Linguistic Regularities in Continuous Space Word Representations. In Proceedings of NAACL HLT, 2013.
- Le, Quoc V., and Tomas Mikolov. "Distributed Representations of Sentences and Documents." ICML. Vol. 14. 2014.
- Tomas Mikolov, Quoc V. Le and Ilya Sutskever. Exploiting Similarities among Languages for Machine Translation.

References

- Levy, Omer, and Yoav Goldberg. "Dependency-Based Word Embeddings." ACL (2). 2014.
- Levy, Omer, Yoav Goldberg, and Ido Dagan. "Improving distributional similarity with lessons learned from word embeddings." *Transactions of the Association for Computational Linguistics* 3 (2015): 211-225.
- Levy, Omer, Yoav Goldberg, and Israel Ramat-Gan. "Linguistic Regularities in Sparse and Explicit Word Representations." CoNLL. 2014.
- L.J.P. van der Maaten and G.E. Hinton. **Visualizing High-Dimensional Data Using t-SNE**. *Journal of Machine Learning Research* 9(Nov):2579-2605, 2008.
- Rajpurkar, Pranav, et al. "SQuAD: 100,000+ Questions for Machine Comprehension of Text." *arXiv preprint arXiv:1606.05250* (2016).

References

- Jolliffe, Ian. *Principal component analysis*. John Wiley & Sons, Ltd, 2002.
- Blei, David M., and John D. Lafferty. "Topic models." *Text mining: classification, clustering, and applications* 10.71 (2009): 34.
- Severyn, Aliaksei, and Alessandro Moschitti. "Learning to rank short text pairs with convolutional deep neural networks." *Proceedings of the 38th International ACM SIGIR Conference on Research and Development in Information Retrieval*. ACM, 2015.